

Persistent Memory Hackathon and Workshop

NVMW '19

March 10, 2019

Contributors:

Jim Fister (SNIA)

Stephen Bates (Eideticom)

Andy Rudoff (Intel)

<https://github.com/pmemhackathon/2019-03-10>

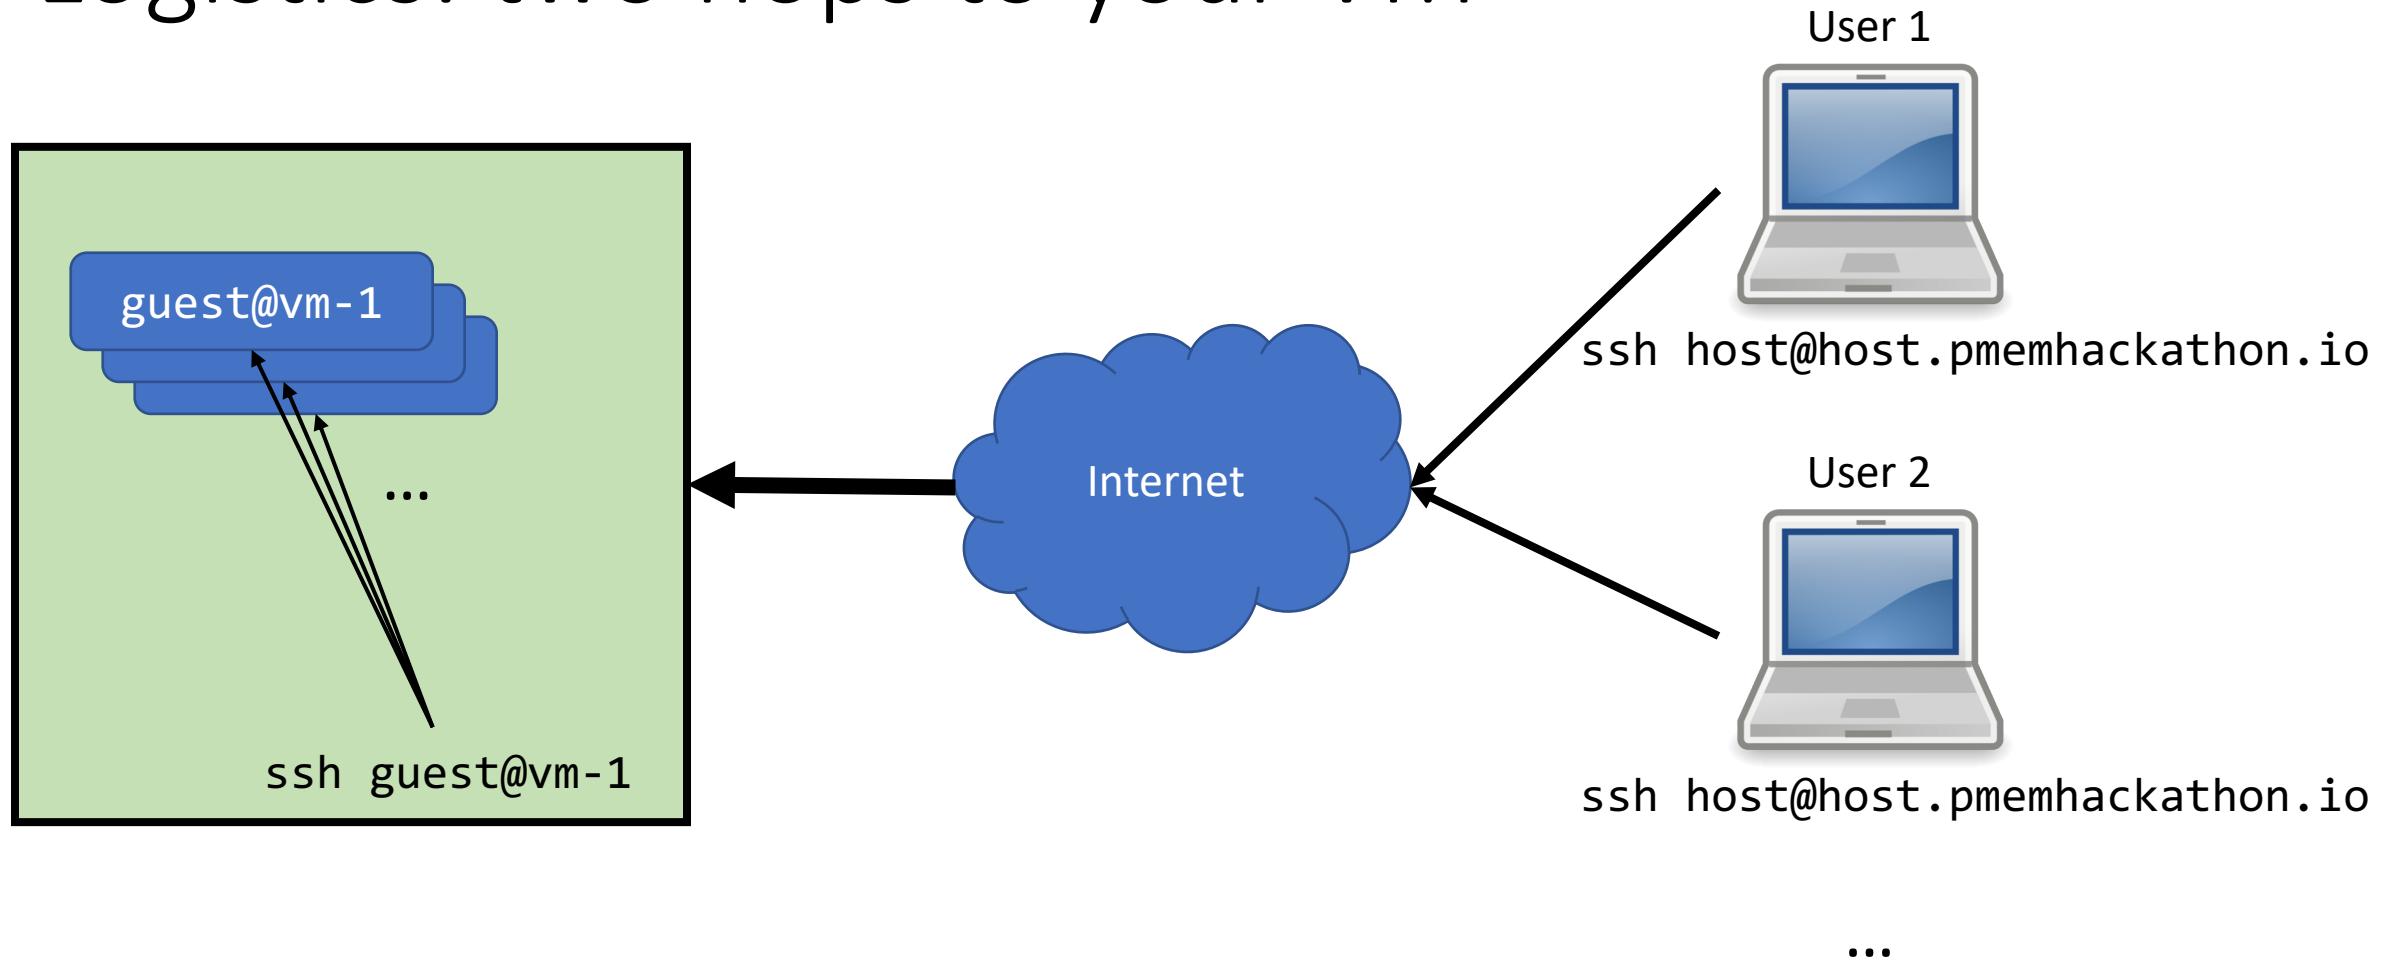
Agenda

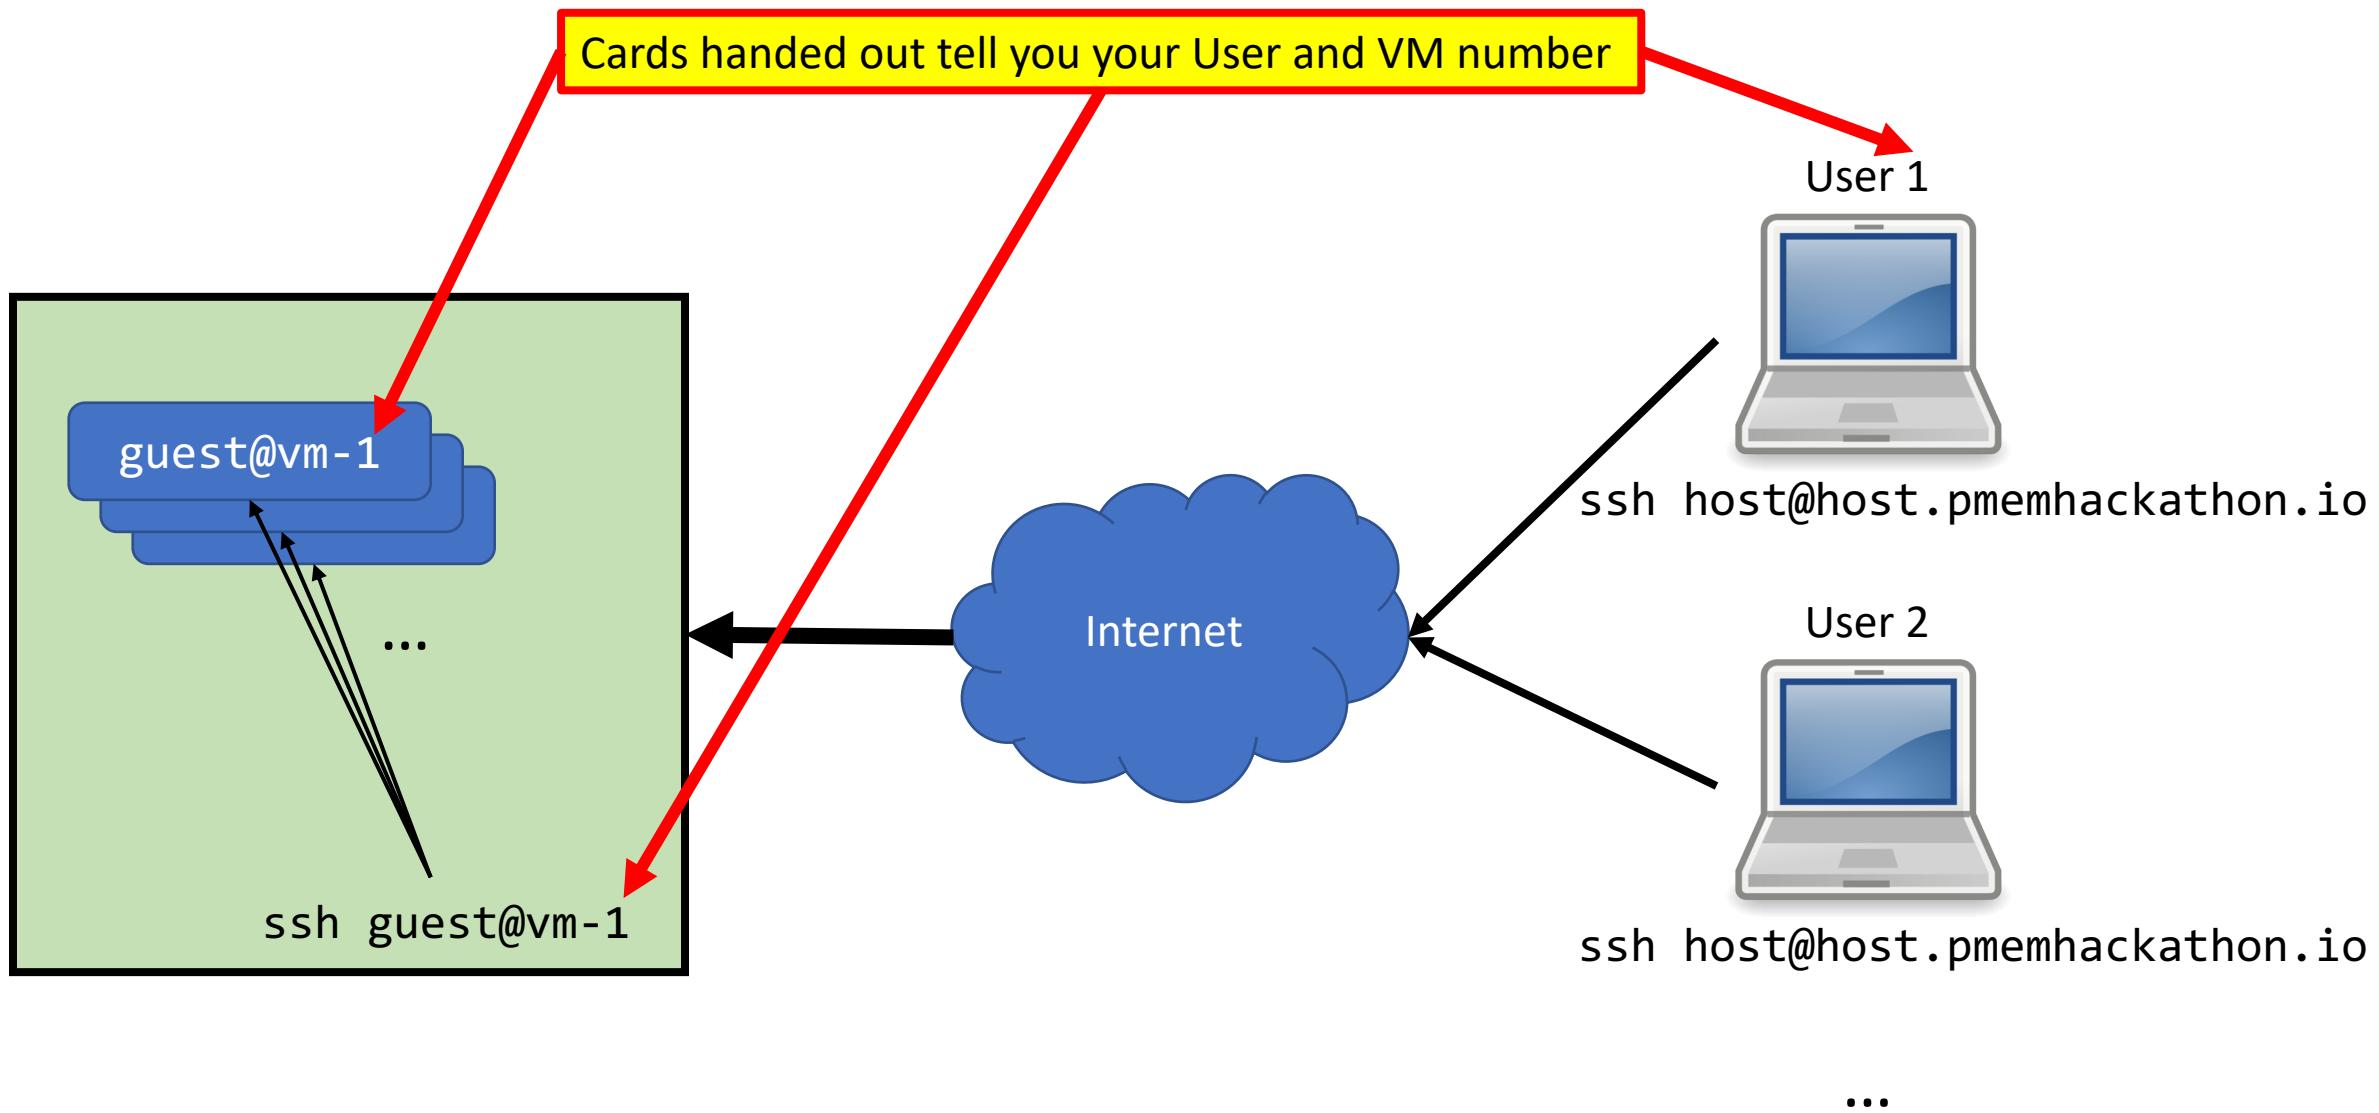
- Logistics
 - How to login to your VM & get it ready
- Persistent Memory Platform Support
 - Platform level support
 - Checking out your kernel
 - Finding and configuring your pmem
- Persistent Memory Programming
 - The basics: storage APIs and memory-mapped files
 - An application's responsibilities when using pmem
 - Installing libraries to help
 - High-level, easy-to-use APIs
 - Lower-level, more flexible (more error-prone) APIs

What Does “Hackathon” Mean To Us?

- Main goal is to show you how to find, configure, and program pmem
 - All slides are in the GitHub repo
 - All shell commands we type are in the GitHub repo
 - You probably don't need to write them down
 - You probably don't even need to type many of them, just cut & paste into the shell
 - Go to <https://github.com/pmemhackathon/2019-03-10> to see today's repo
 - But in a minute, we'll demonstrate cloning the repo to your VM
- Mostly we will show you how to install stuff and get you going
 - After installing samples, try them out, or write your own
 - We'll walk through some for everyone, then will walk around & help you

Logistics: two hops to your VM





VM Basics

- Set your hostname
 - `sudo sh -c 'echo vm-1 >/etc/hostname'`
- Make sure your SW is up to date
 - `sudo apt update && sudo apt upgrade`
 - `sudo reboot`
- NOTE: Today's Hackathon is using **emulated pmem!**
 - (but some demos will be on the real stuff)

Make a local clone of the hackathon repo

```
$ cd  
$ git clone https://github.com/pmemhackathon/2019-03-10  
Cloning into '2019-03-10'...  
remote: Enumerating objects: 14, done.  
remote: Counting objects: 100% (14/14), done.  
remote: Compressing objects: 100% (13/13), done.  
remote: Total 14 (delta 1), reused 14 (delta 1), pack-reused 0  
Unpacking objects: 100% (14/14), done.  
$ cd 2019-03-10  
$ more README.txt
```

Most of the shell commands we type during demos are in this README.txt

Does your System Support Persistent Memory?

- Does my platform support persistent memory?
 - Your vendor determines this. Buy a system meant for it.
 - Don't just buy an NVDIMM and plug it into a random system – you need platform support (like BIOS, ADR, power supply). You want validated configurations.

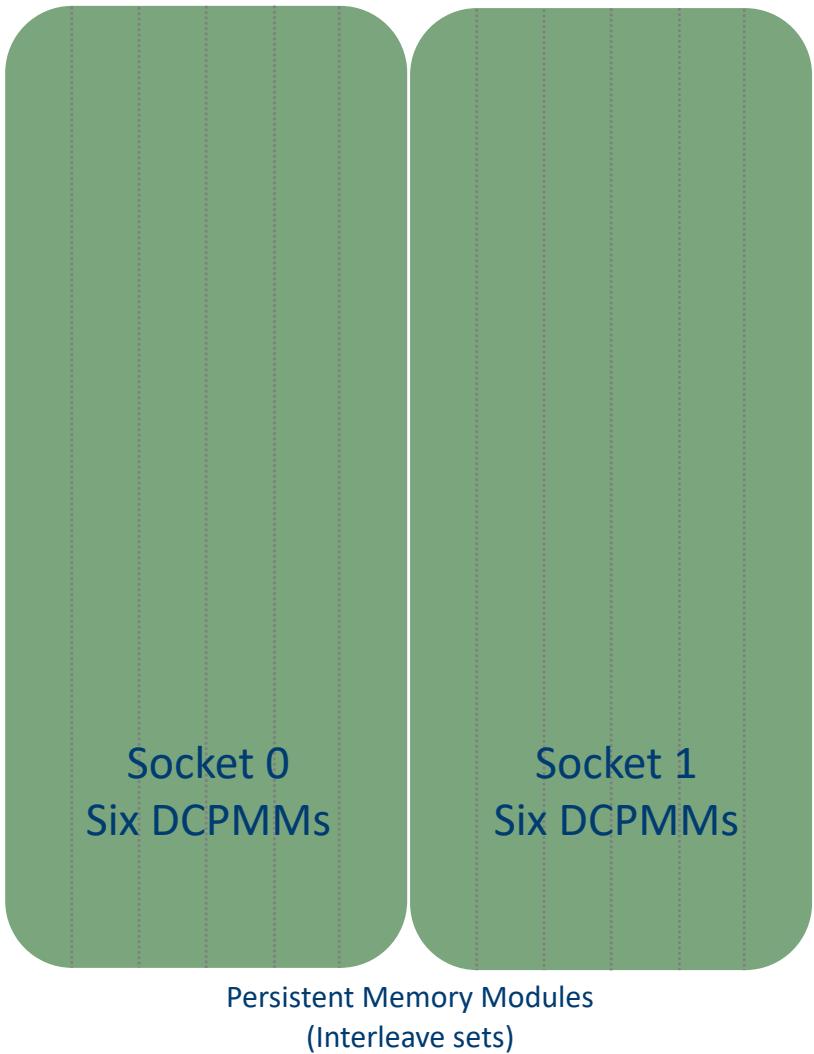
```
ndctl list -BN      # check the "provider" field for ACPI.NFIT
```

- Does my OS support persistent memory?
 - Major OS vendors (and Linux distros) will tell you which version supports it
 - Linux kernel support is enabled in the config file used to build the kernel

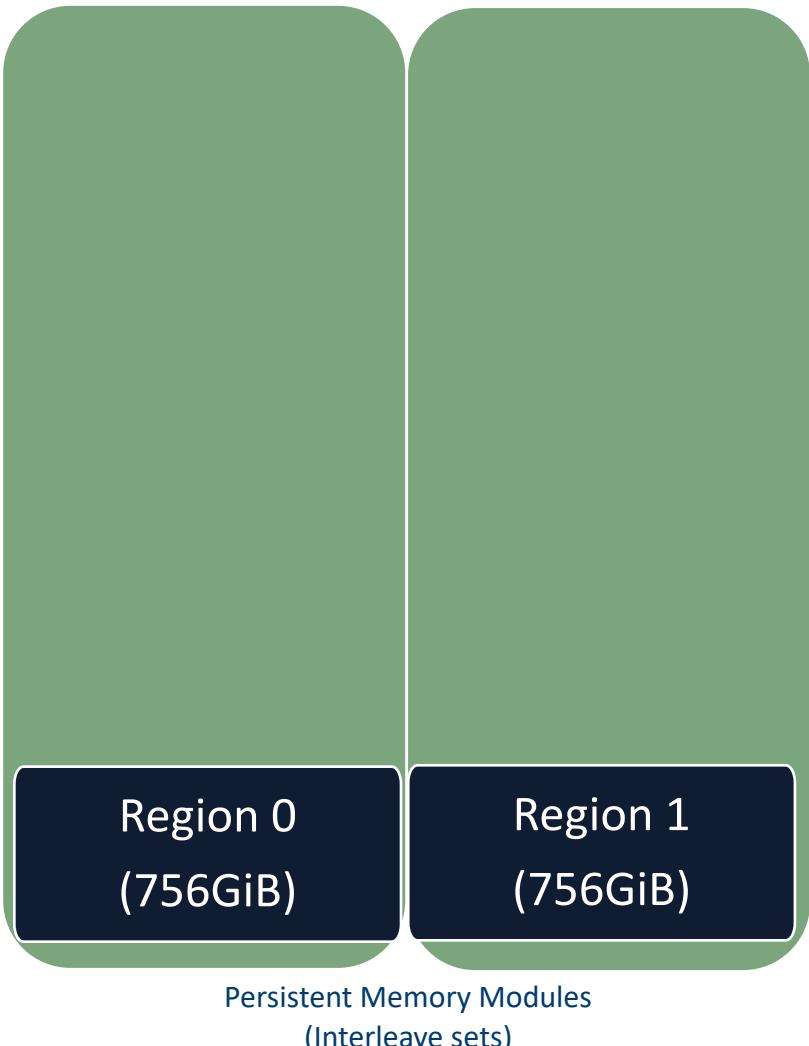
```
uname -r          # see kernel currently running
grep -i pmem /boot/config-`uname -r`
grep -i nvdimm /boot/config-`uname -r`
```

Example: Provisioning Intel® Optane DC Persistent Memory Modules

Hardware

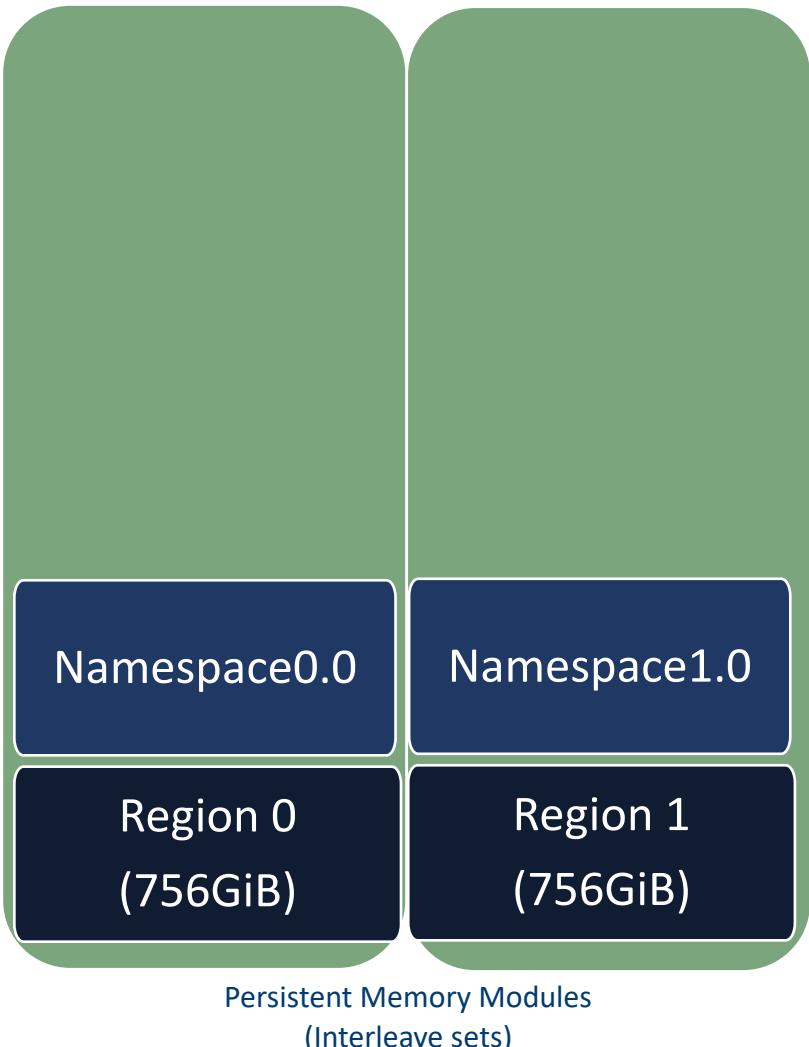


Hardware



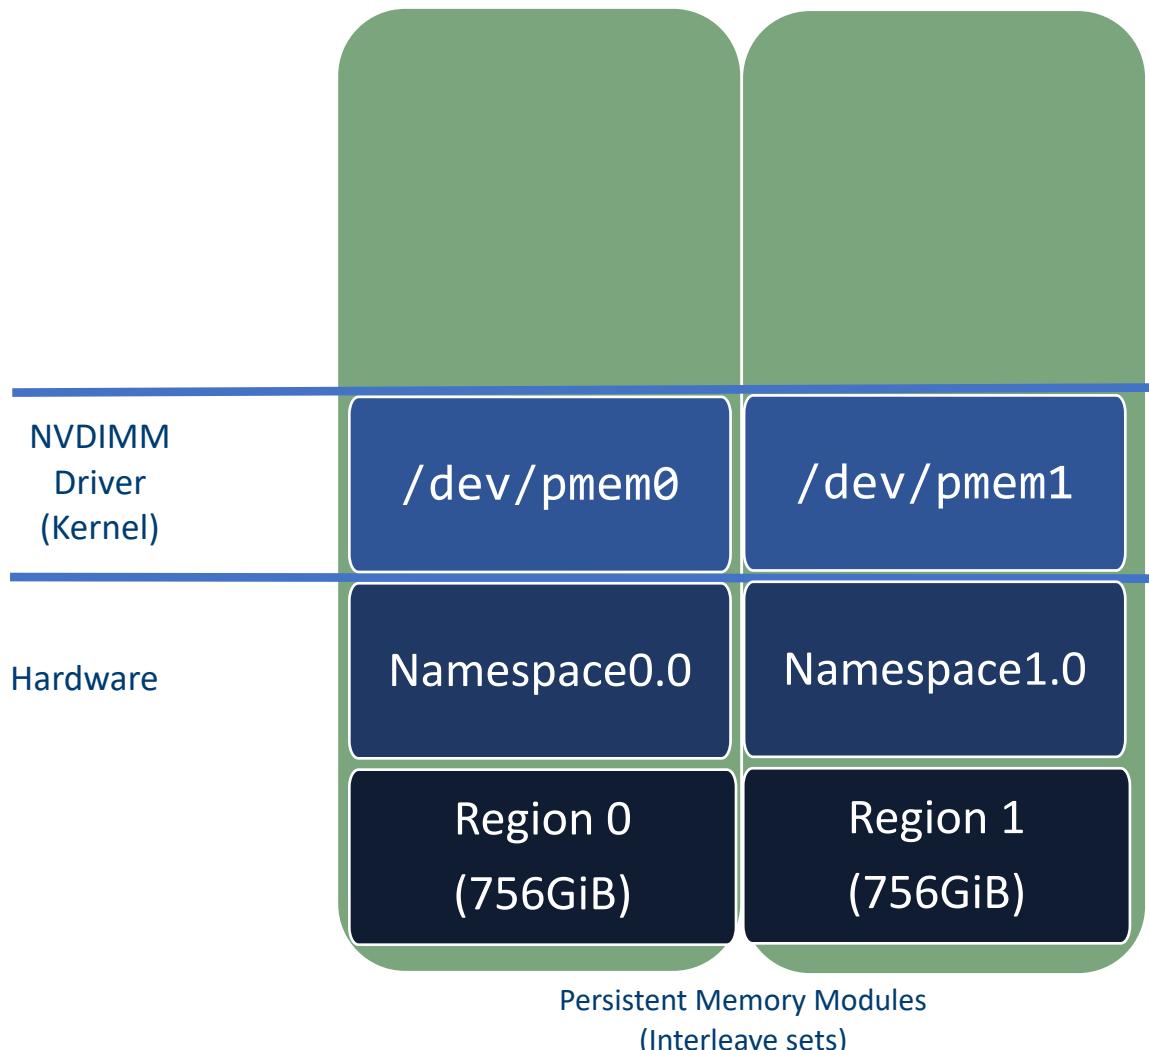
```
# ipmctl create -goal PersistentMemoryType=AppDirect
```

Hardware



```
# ndctl create-namespace  
# ndctl create-namespace
```

```
# ipmctl create -goal PersistentMemoryType=AppDirect
```



```
# ndctl create-namespace
# ndctl create-namespace
```

```
# ipmctl create -goal PersistentMemoryType=AppDirect
```

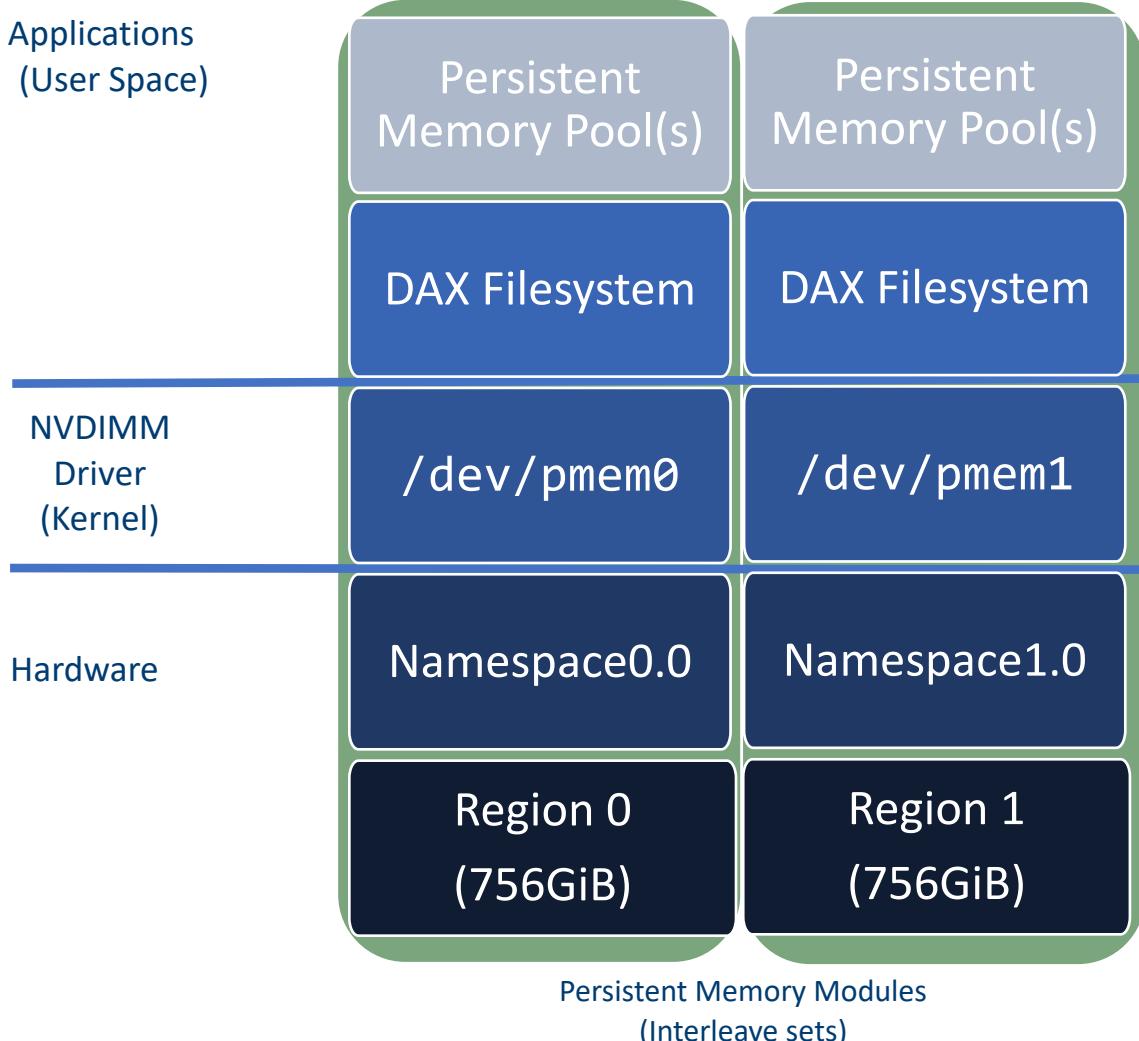
Applications
(User Space)



```
# mkfs.ext4 /dev/pmem0
# mkfs.ext4 /dev/pmem1
# mount -o dax /dev/pmem0 /mnt/pmem-fsdax0
# mount -o dax /dev/pmem1 /mnt/pmem-fsdax1
```

```
# ndctl create-namespace
# ndctl create-namespace
```

```
# ipmctl create -goal PersistentMemoryType=AppDirect
```



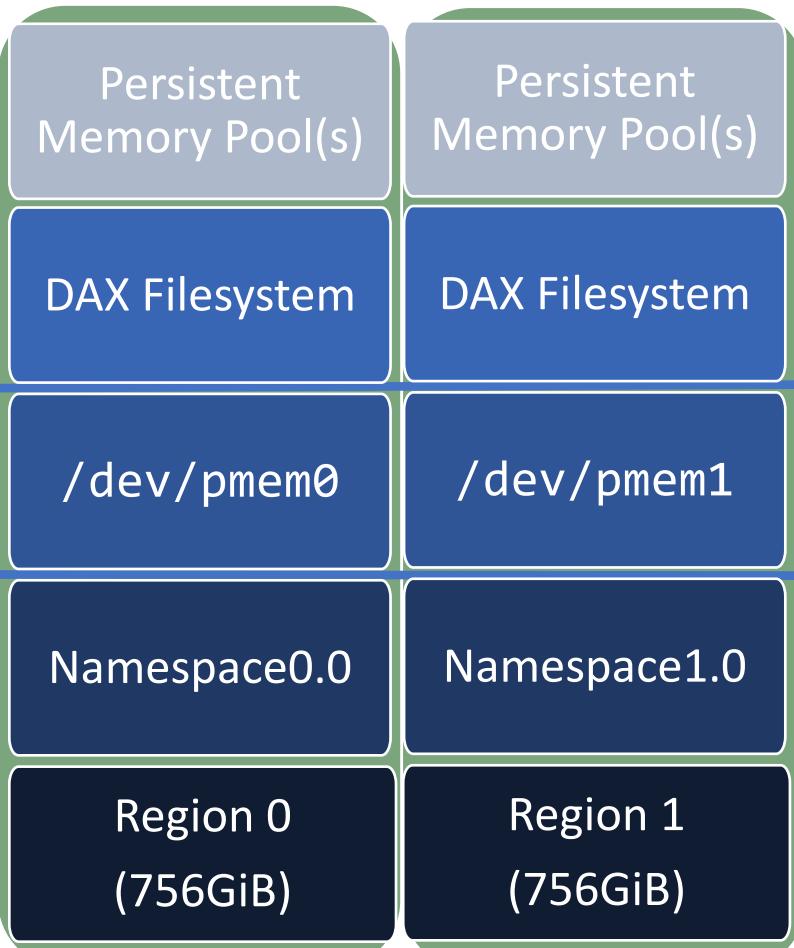
```
$ pmempool create /mnt/pmem-fsdax0/pool0
$ pmempool create /mnt/pmem-fsdax1/pool1
```

```
# mkfs.ext4 /dev/pmem0
# mkfs.ext4 /dev/pmem1
# mount -o dax /dev/pmem0 /mnt/pmem-fsdax0
# mount -o dax /dev/pmem1 /mnt/pmem-fsdax1
```

```
# ndctl create-namespace
# ndctl create-namespace
```

```
# ipmctl create -goal PersistentMemoryType=AppDirect
```

Applications
(User Space)



```
$ pmempool create /mnt/pmem-fsdax0/pool0  
$ pmempool create /mnt/pmem-fsdax1/pool1
```

```
# mkfs.ext4 /dev/pmem0  
# mkfs.ext4 /dev/pmem1  
# mount -o dax /dev/pmem0 /mnt/pmem-fsdax0  
# mount -o dax /dev/pmem1 /mnt/pmem-fsdax1
```

```
# ndctl create-namespace  
# ndctl create-namespace
```

Vendor Neutral

```
# ipmctl create -goal PersistentMemoryType=AppDirect
```

Vendor Specific

In your VM...

```
$ sudo ndctl list -u
$ sudo ndctl create-namespace -f -e namespace0.0 --mode fsdax

$ ls -l /dev/pmem*

$ sudo mkfs.ext4 /dev/pmem0

$ sudo mkdir /mnt/pmem-fsdax
$ sudo mount -o dax /dev/pmem0 /mnt/pmem-fsdax

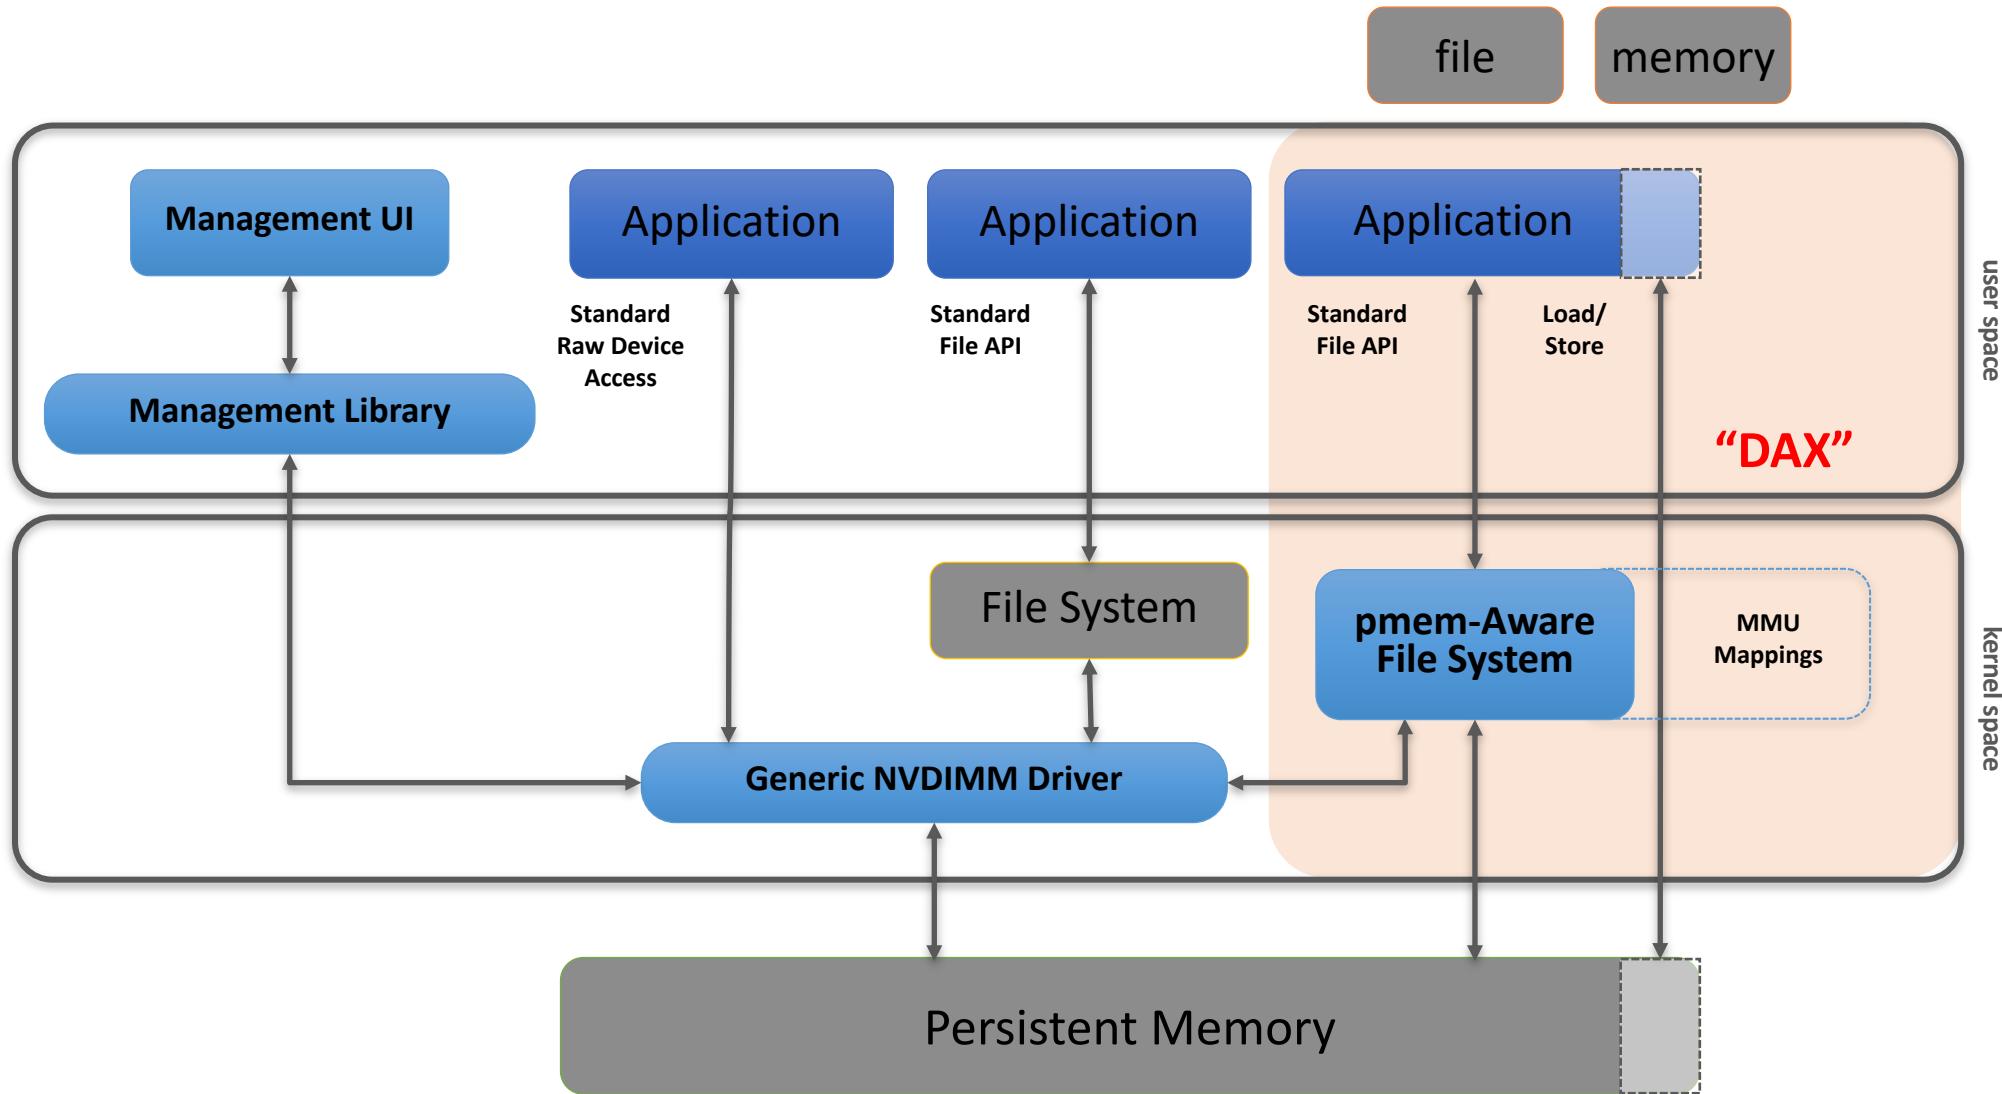
$ sudo chmod 777 /mnt/pmem-fsdax    # open up perms for this hackathon

$ df -h
... other file-related stuff works as expected...
```

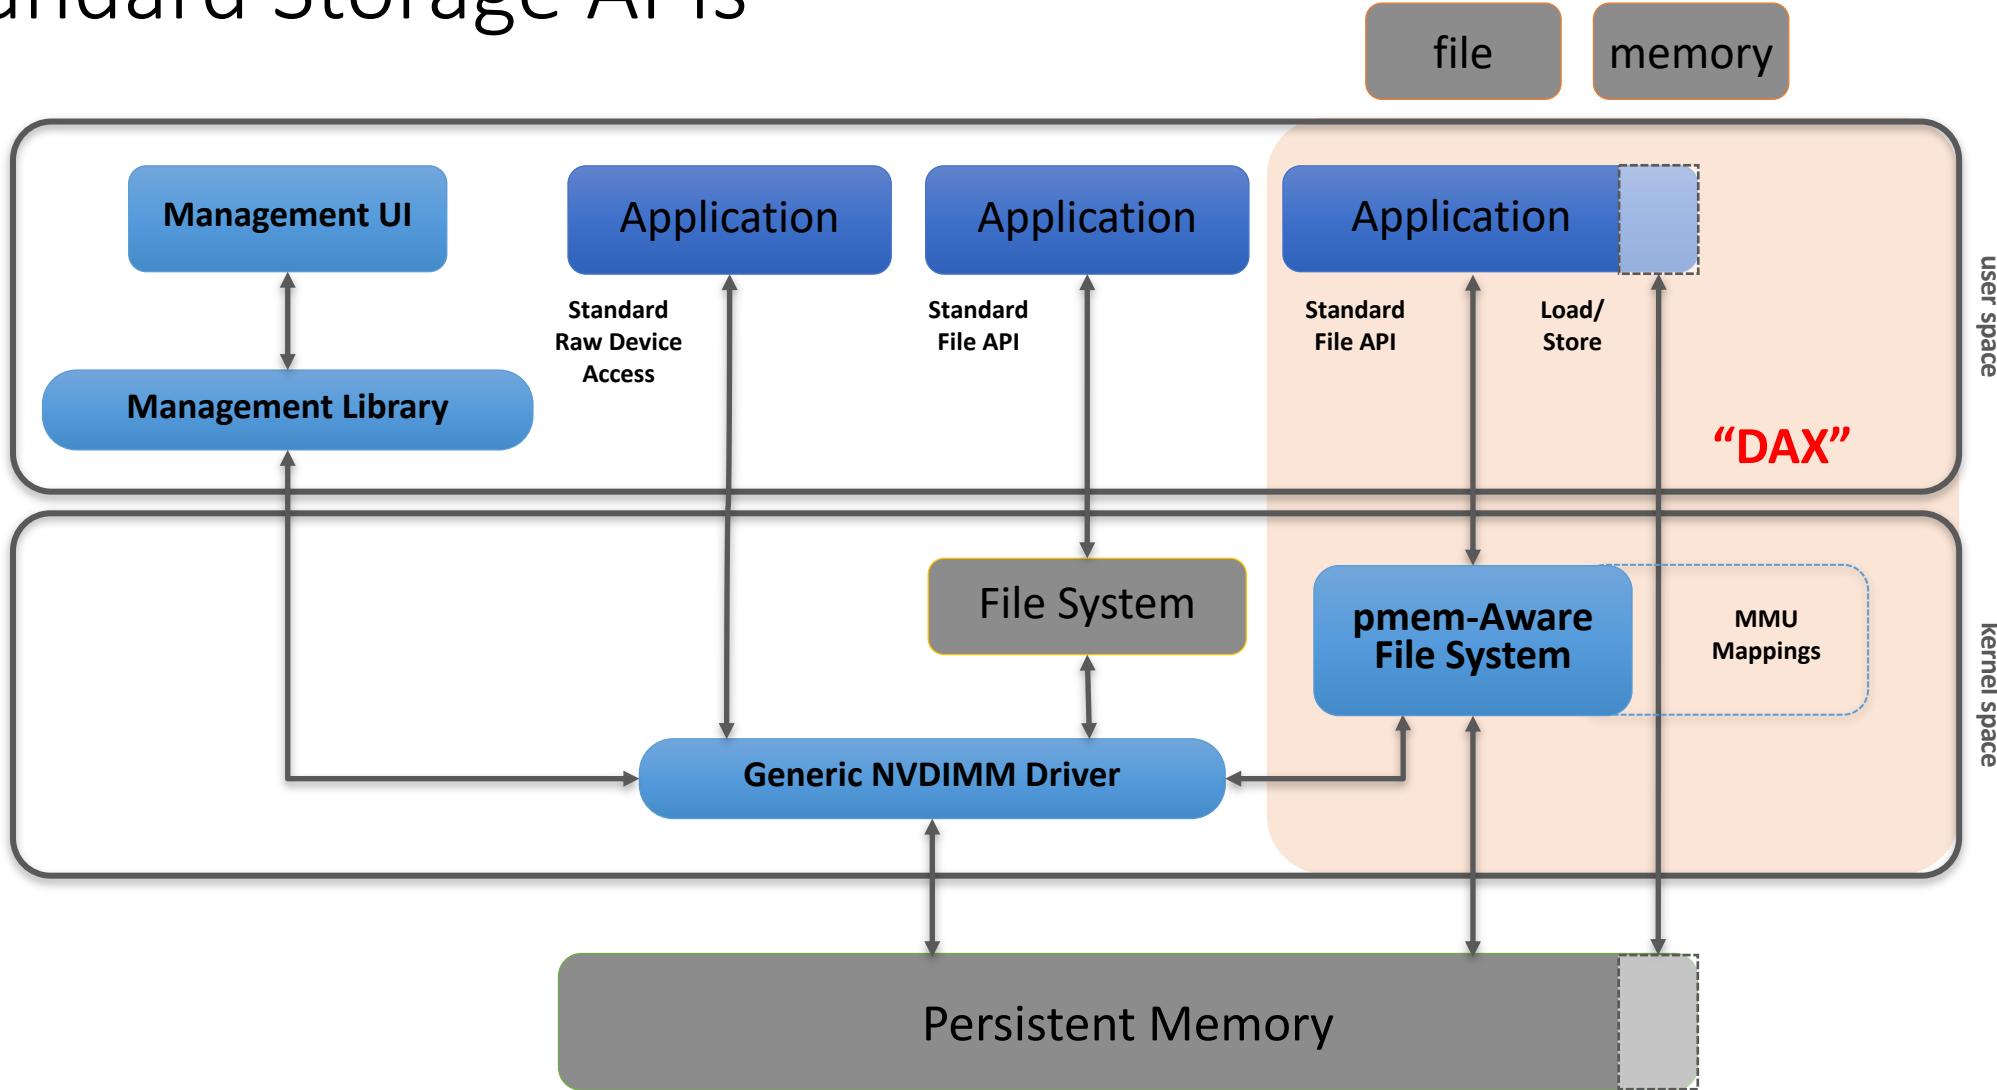
Essential Programming Background

- Lots of ways to use pmem with existing programs
 - Storage APIs
 - Libraries or kernels using pmem transparently
 - Memory Mode
- This hackathon doesn't cover the above (too easy!)
 - We assume you want direct access to pmem
 - We show code, but also concepts
 - There are lots of paths you can take, these are just examples

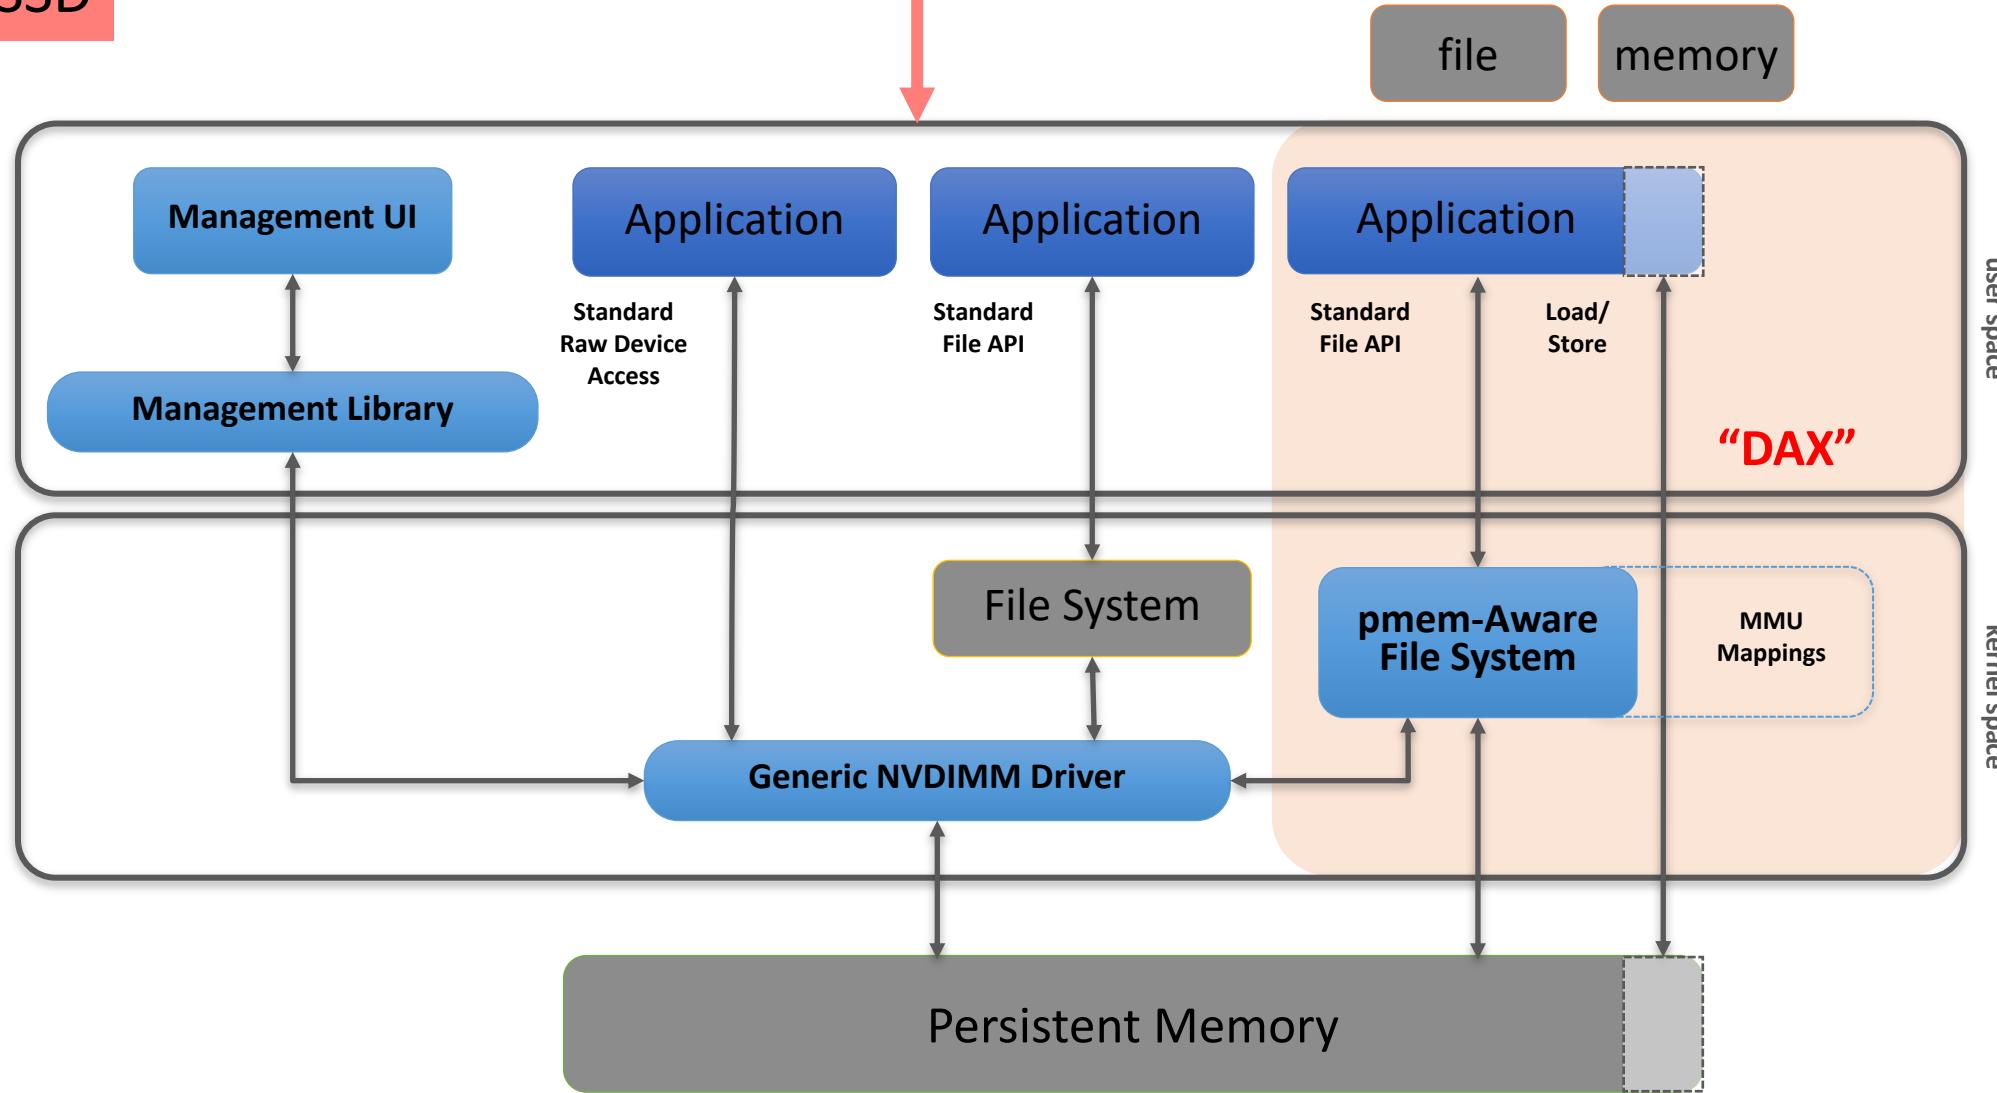
The SNIA NVM Programming Model



Don't Forget: The NVM Programming Model Starts With Standard Storage APIs

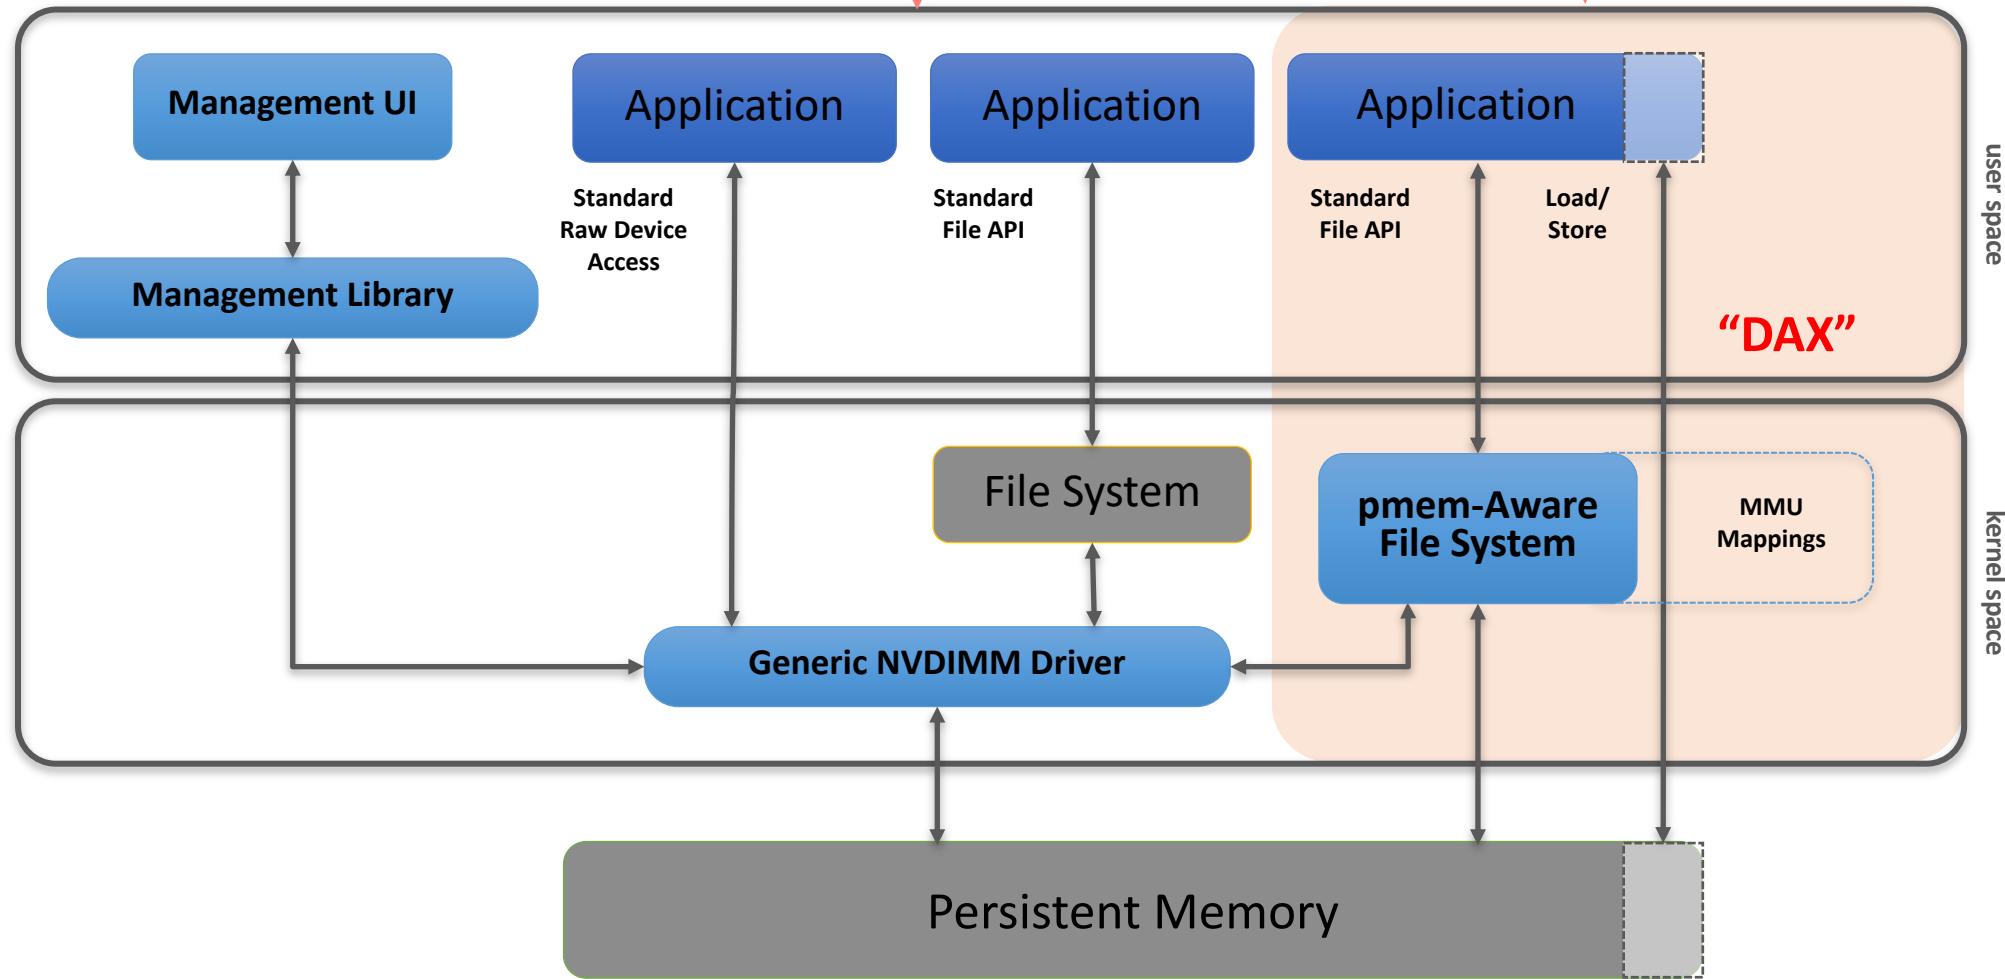


Use PM
Like an SSD



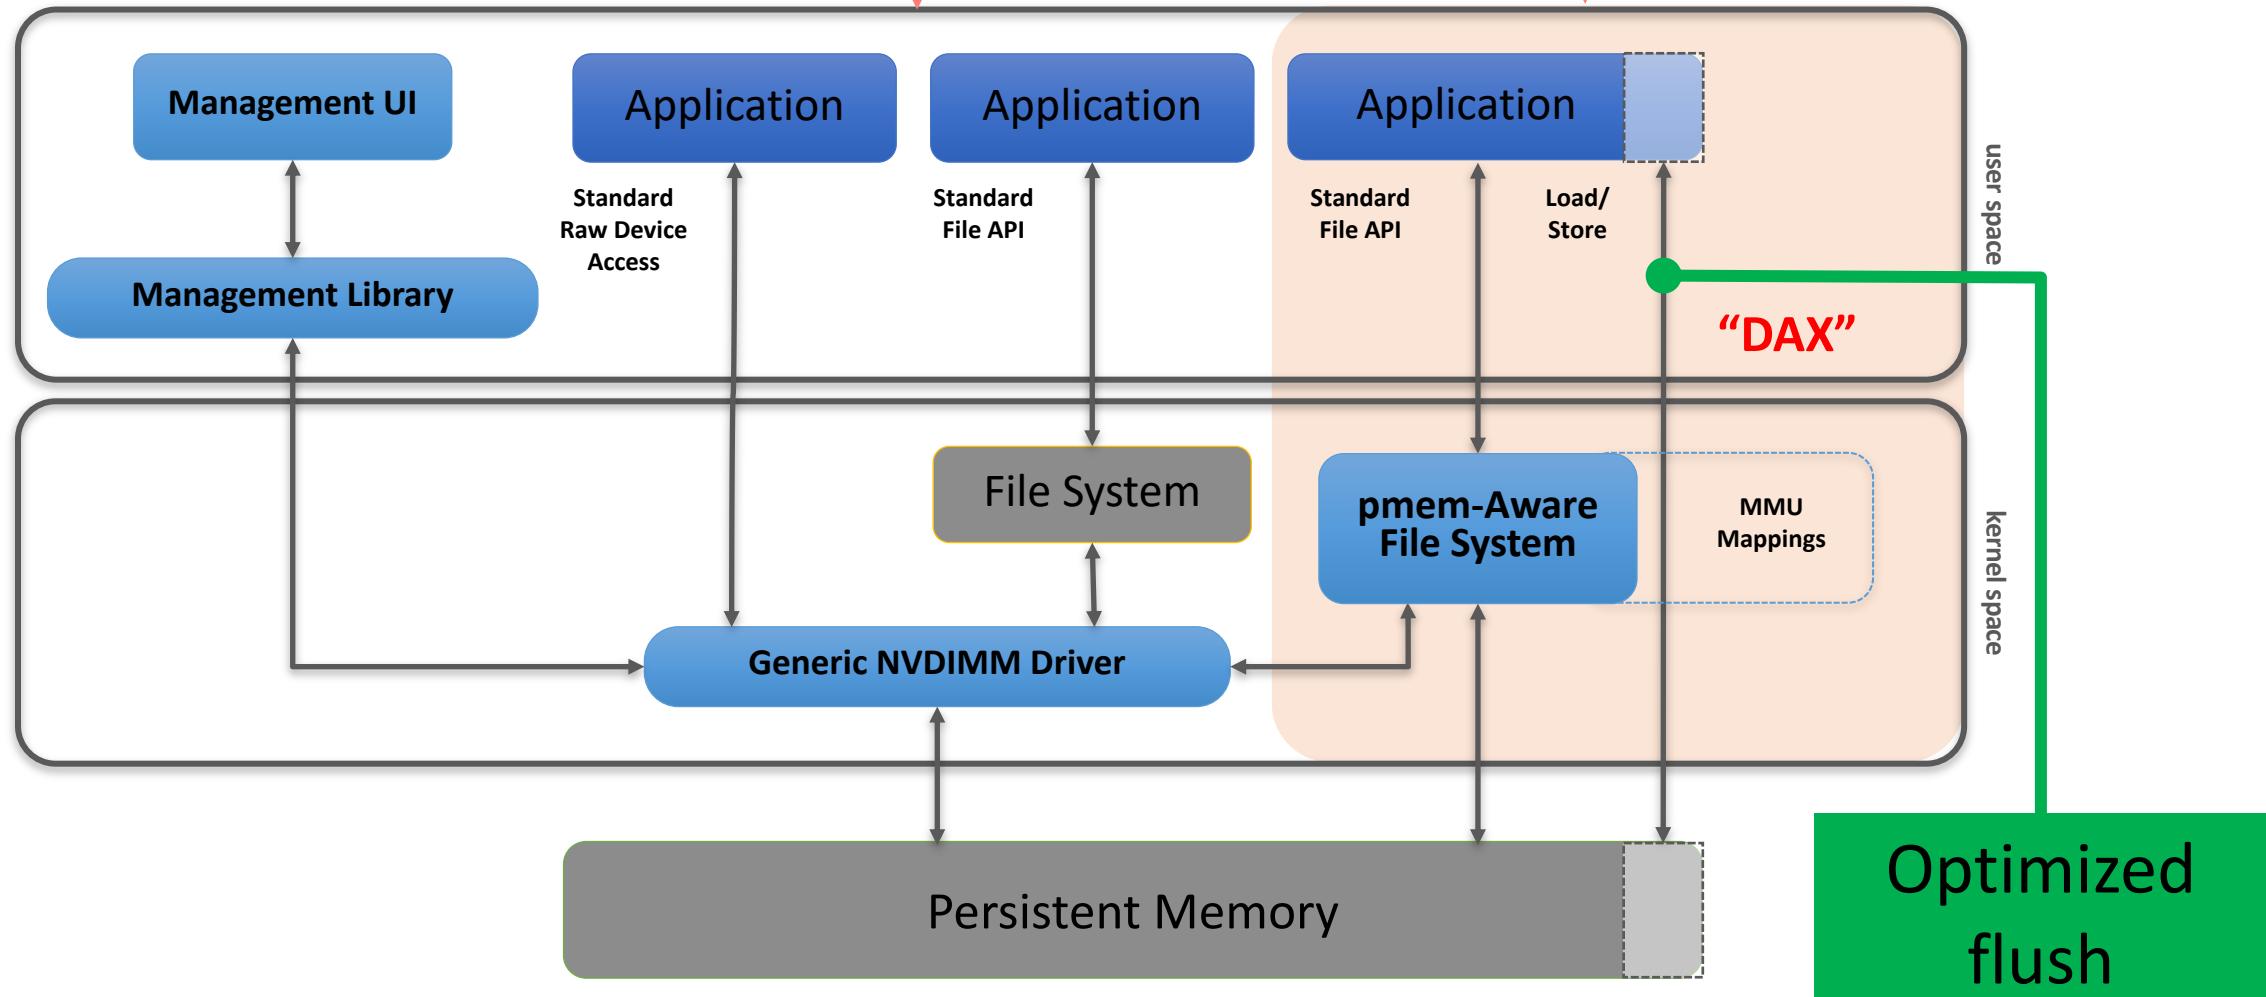
Use PM
Like an SSD

Use PM
Like an SSD
(no page cache)



Use PM
Like an SSD

Use PM
Like an SSD
(no page cache)



Programming Examples For This Hackathon

We stick to the DAX path

- RAW Access to pmem
 - mmap() -- you get a pointer to pmem, the rest is up to you
 - Only 8-byte stores are powerfail atomic
- libmemkind
 - volatile use of pmem
- Pointers to high-level persistent abstractions available
 - libpmemkv
 - C++ vector
- libpmem
 - One step above RAW access, still only 8-byte stores are powerfail atomic
 - mmap(), memcpy() helper functions, optimized flush functions
- libpmemblk
 - Very simple transactional library, read/write fixed sized block only
- libpmemobj
 - **General-purpose allocations, transactions, atomics (series of examples)**

Programming Examples For This Hackathon

We stick to the DAX path

- RAW Access to pmem
 - mmap() -- you get a pointer to pmem, the rest is up to you
 - Only 8-byte stores are powerfail atomic
- libmemkind
 - volatile use of pmem
- Pointers to high-level persistent abstractions available
 - libpmemkv
 - C++ vector
- libpmem
 - One step above RAW access, still only 8-byte stores are powerfail atomic
 - mmap(), memcpy() helper functions, optimized flush functions
- libpmemblk
 - Very simple transactional library, read/write fixed sized block only
- libpmemobj
 - General-purpose allocations, transactions, atomics (series of examples)

Deciding what to use...

Are standard storage APIs good enough?

Do you need persistence?

High-level abstraction available?

Build your own

Resources

- PMDK Resources:
 - Home: <https://pmem.io>
 - PMDK: <https://pmem.io/pmdk>
 - PMDK Source Code : <https://github.com/pmem/PMDK>
 - Google Group: <https://groups.google.com/forum/#!forum/pmem>
 - Intel Developer Zone: <https://software.intel.com/persistent-memory>
 - Memkind: <https://github.com/memkind/memkind> (see memkind_pmem(3))
 - libpmemkv: <https://github.com/pmem/pmemkv>
- NDCTL: <https://pmem.io/ndctl>
- SNIA NVM Programming Model:
https://www.snia.org/tech_activities/standards/curr_standards/npm
- Getting Started Guides: <https://docs.pmem.io>

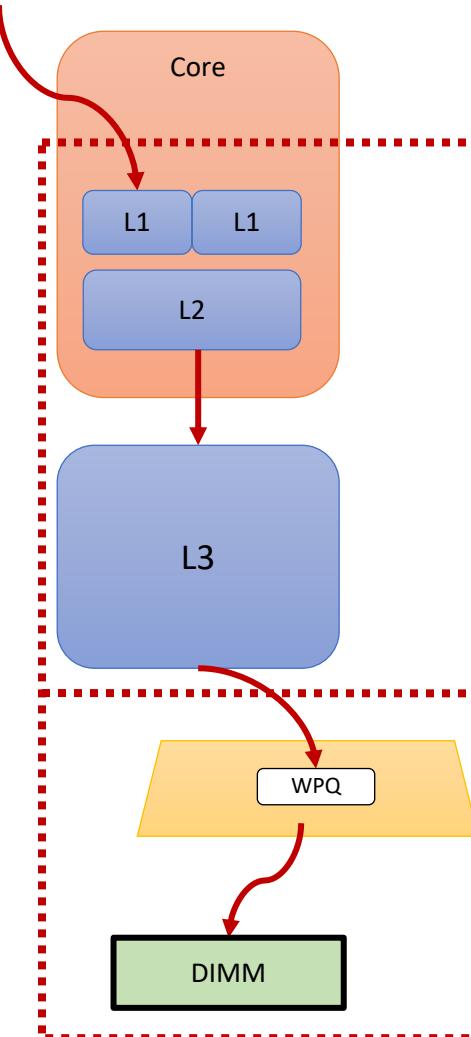
A Programmer's View (mapped files)

```
fd = open("/my/file", O_RDWR);  
...  
base = mmap(NULL, filesize,  
            PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);  
close(fd);  
...  
base[100] = 'X';  
strcpy(base, "hello there");  
*structp = *base_structp;  
...
```

“Load/Store”

How the Hardware Works

MOV



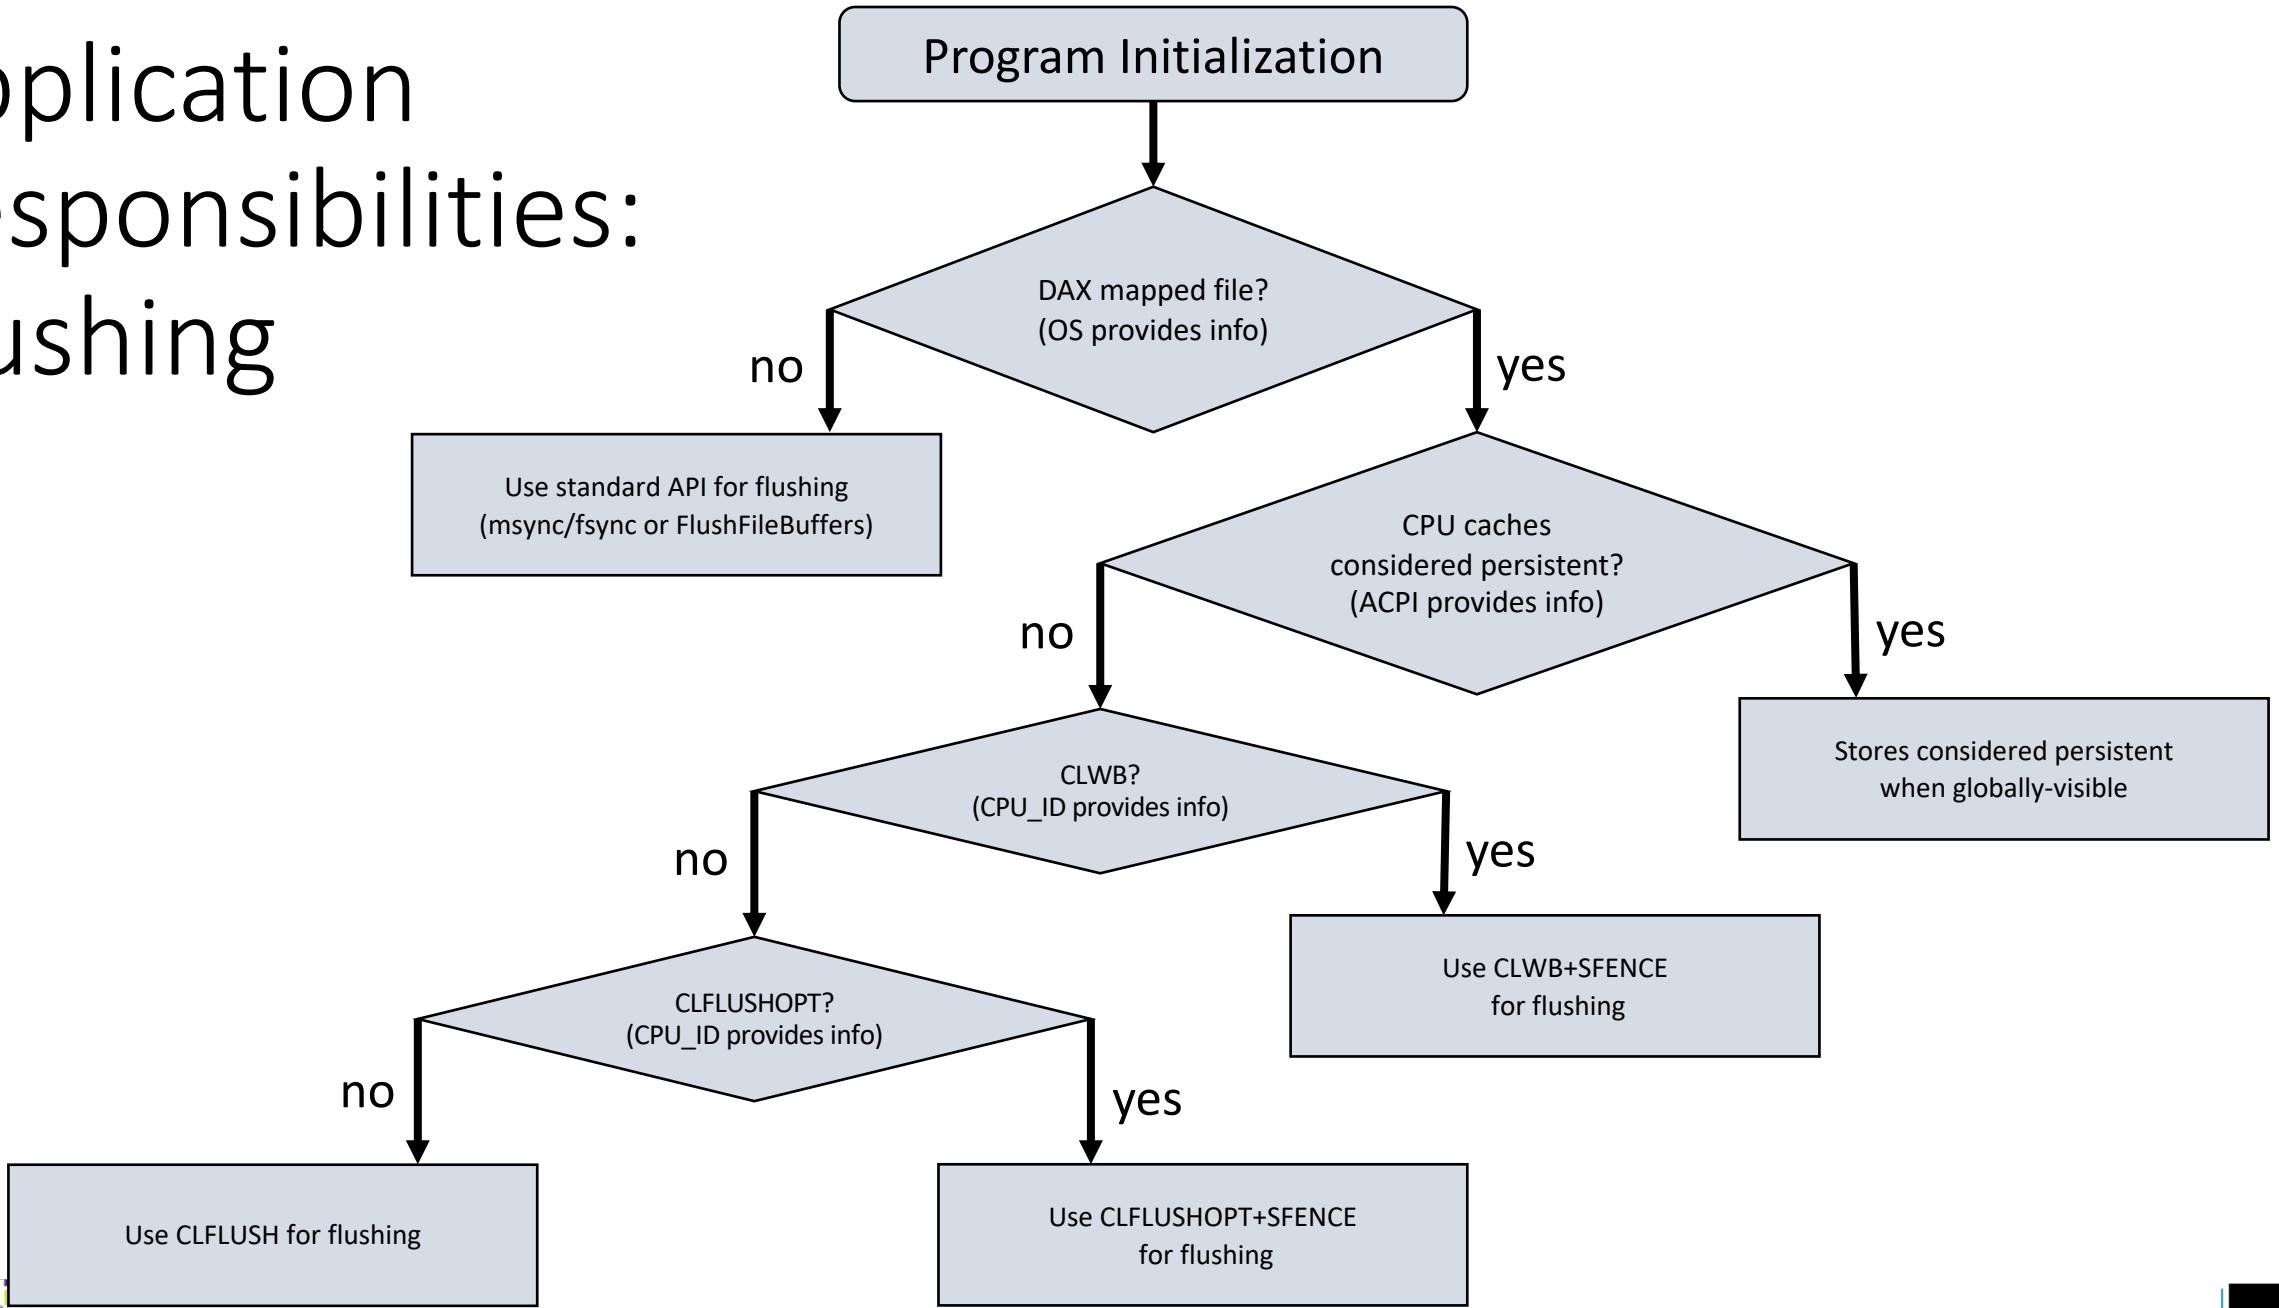
CLWB + fence
-or-
CLFLUSHOPT + fence
-or-
CLFLUSH
-or-
NT stores + fence
-or-
WBINVD (kernel only)

Custom
Power fail protected domain
indicated by ACPI property:
CPU Cache Hierarchy

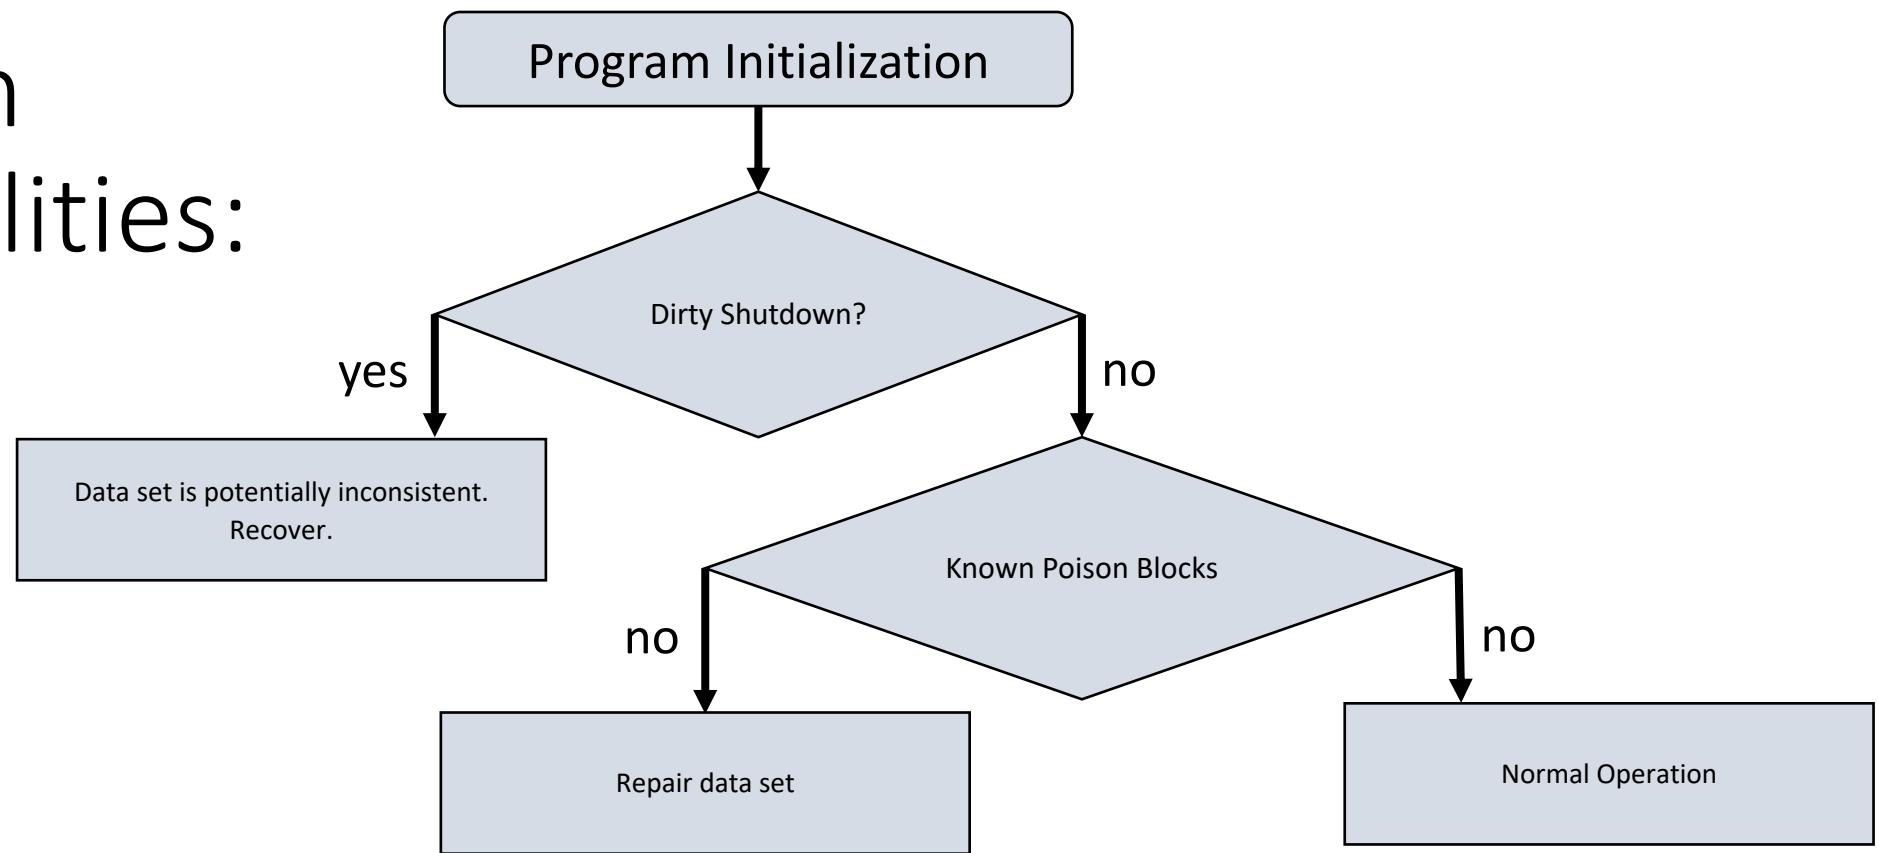
Minimum Required
Power fail protected domain:
Memory subsystem

Not shown:
MCA
ADR Failure Detection

Application Responsibilities: Flushing



Application Responsibilities: Recovery



Application Responsibilities: Consistency

```
open(...);  
  
mmap(...);  
  
strcpy(pmemp, "Hello, World!");  
  
msync(...);
```

← Crash

Result

1. "\0\0\0\0\0\0\0\0\0\0..."
2. "Hello, \0\0\0\0\0\0..."
3. "\0\0\0\0\0\0\0\0\0orld!\0"
4. "Hello, \0\0\0\0\0\0\0\0\0"
5. "Hello, World!\0"

Application Responsibilities: Consistency

```
open(...);  
mmap(...);  
strcpy(pmemp, "Hello, World!");  
pmem_persist(pmemp, 14);
```

`pmem_persist()` may be faster,
but is still **not** transactional

← Crash

Result

1. "\0\0\0\0\0\0\0\0\0\0..."
2. "Hello, \0\0\0\0\0\0..."
3. "\0\0\0\0\0\0\0\0\0orld!\0"
4. "Hello, \0\0\0\0\0\0\0\0\0"
5. "Hello, World!\0"

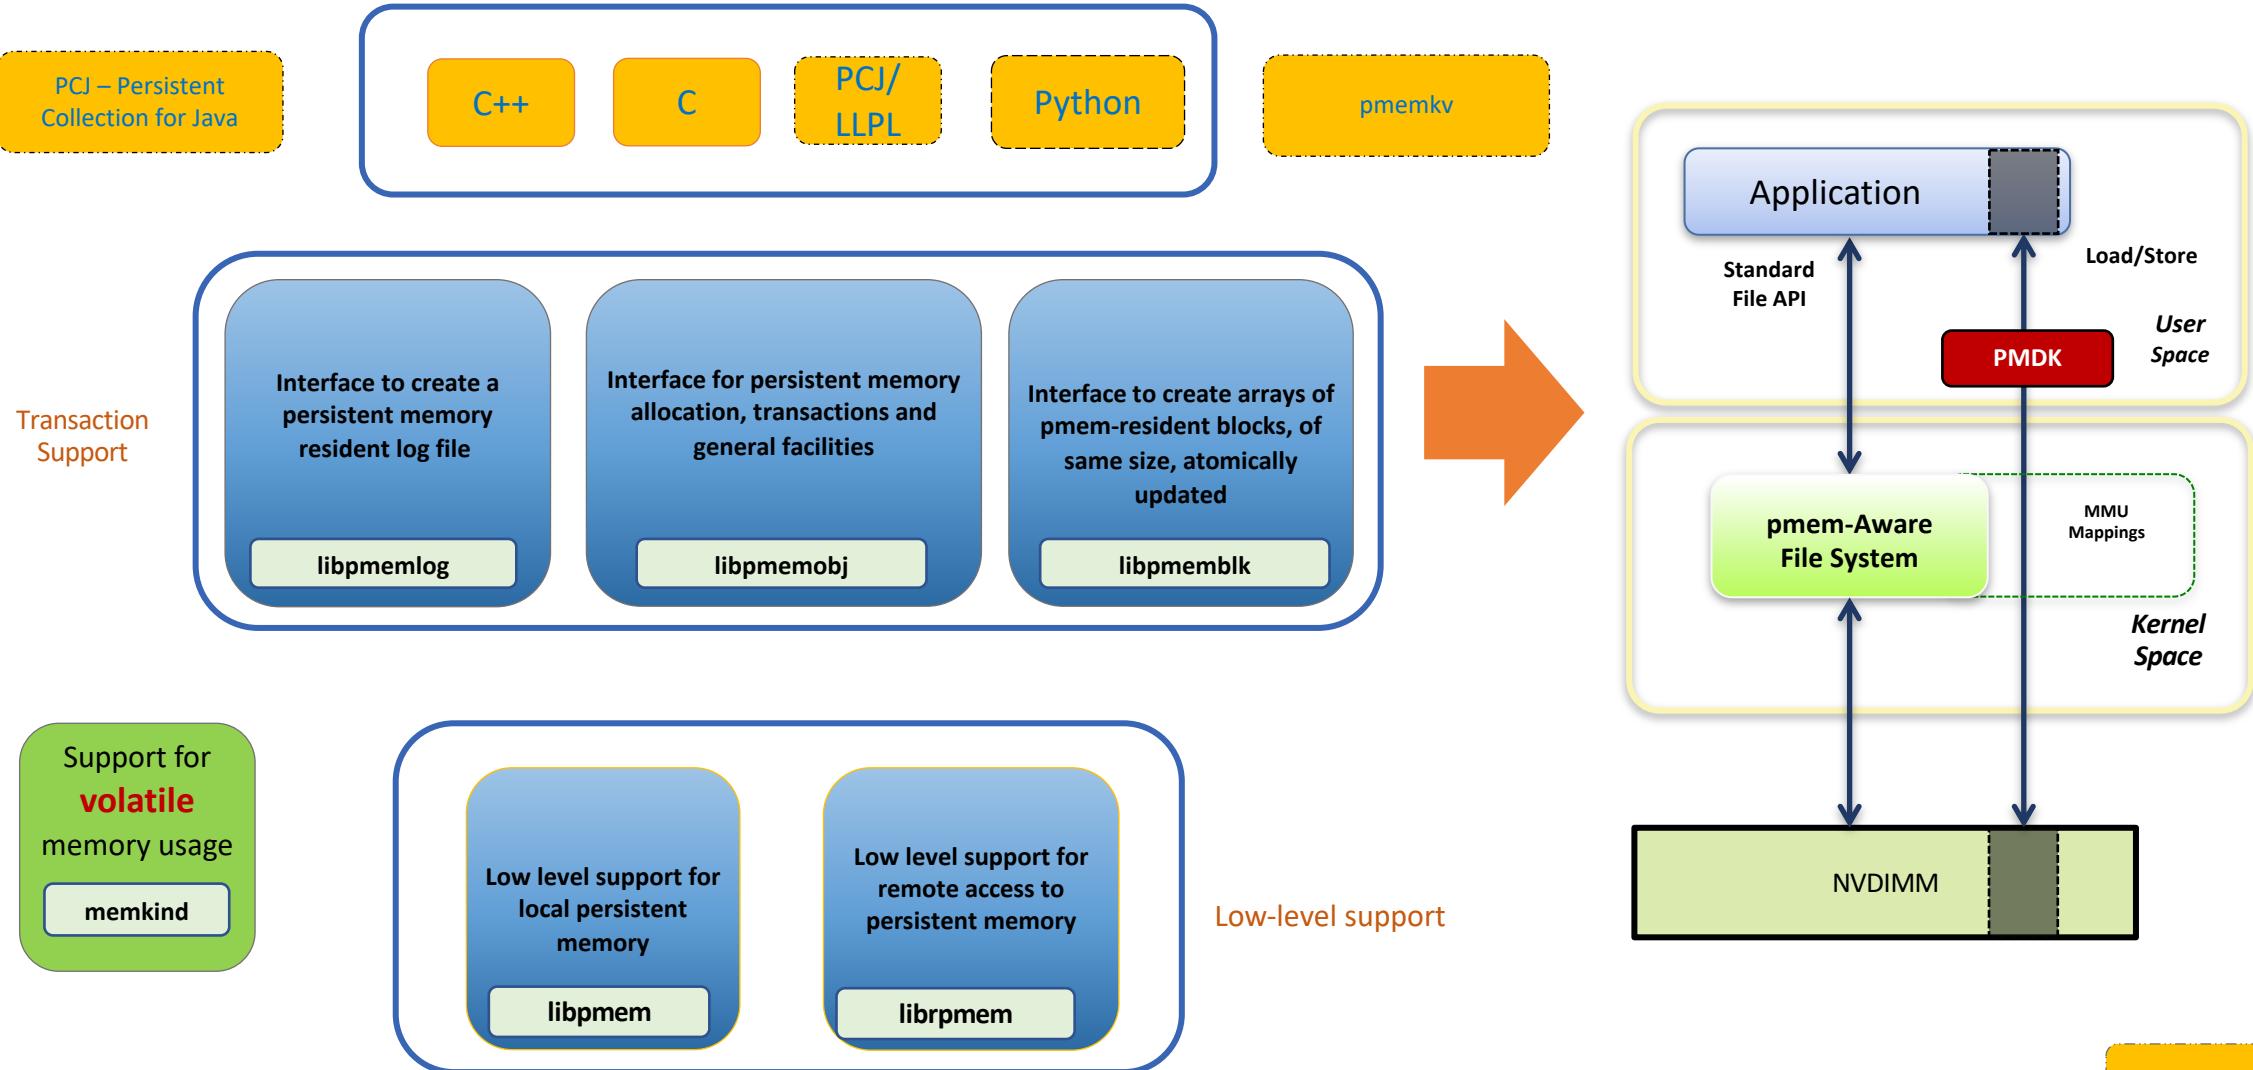
The Persistent Memory Development Kit

PMDK <http://pmem.io>

- PMDK is a collection of libraries
 - Developers pull only what they need
 - Low level programming support
 - Transaction APIs
 - Fully validated
 - Performance tuned.
- Open Source & Product neutral



PMDK Libraries



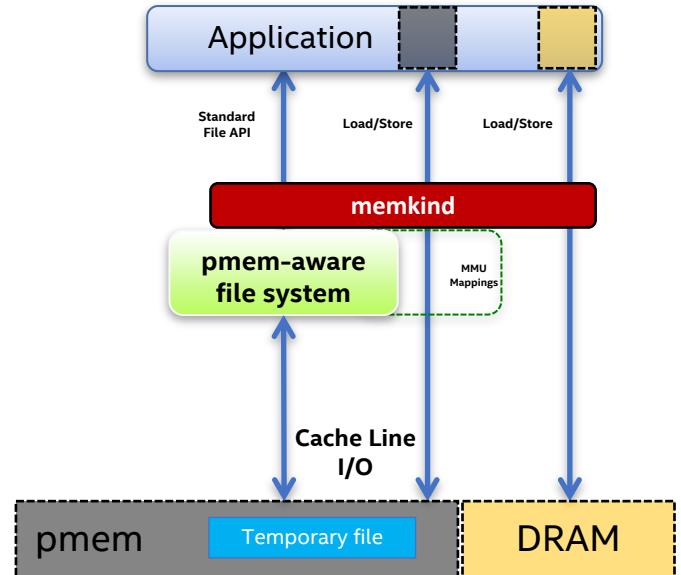
Installing libraries

- Many libraries discussed during this hackathon are upstream
- Product & libraries are fairly new
 - Lots of recent enhancements, especially around performance
- We show installing the libraries by cloning & building the source
 - This gets you the latest & greatest
 - Also gets you in-tree tools & examples
 - Encourages you to submit issues to GitHub
- For products, we would expect apps to check dependencies
 - Ensuring the libraries installed on the system were validated for that app

libmemkind: Volatile use of pmem

(Not for NVDIMM-N, but popular use for Optane DCPMM)

- When to use:
 - Application wants capacity, cheaper/larger than DRAM
 - Application doesn't need persistence for this data
- Normal C malloc library doesn't support multiple pools of memory
- libmemkind does
 - For pmem it creates a temporary file,



```
memkind_create_pmem(const char *dir, size_t max_size, memkind_t *kind)  
memkind_malloc(memkind_t kind, size_t size)  
memkind_calloc(memkind_t kind, size_t num, size_t size)  
memkind_realloc(memkind_t kind, void *ptr, size_t size)  
memkind_free(memkind_t kind, void *ptr)
```

Using High-Level Abstractions

- libpmemkv
 - <https://github.com/pmem/pmemkv>
 - Still “experimental” but product-level validation in progress
- C++ containers
 - <http://pmem.io/2019/02/20/cpp-vector.html>
 - More containers under development

libpmem examples

Source: <https://github.com/pmem/pmdk/tree/master/src/examples/libpmem>

```
/*
 * simple_copy.c -- show how to use pmem_memcpy_persist()
 *
 * usage: simple_copy src-file dst-file
 *
 * Reads 4k from src-file and writes it to dst-file.
 */
```

```
/* create a pmem file and memory map it */
if ((pmemaddr = pmem_map_file(argv[2], BUF_LEN,
                               PMEM_FILE_CREATE|PMEM_FILE_EXCL,
                               0666, &mapped_len, &is_pmem)) == NULL) {
    perror("pmem_map_file");
    exit(1);
}
```

Using is_pmem

```
if (is_pmem) {  
    pmem_memcpy_persist(pmemaddr, buf, cc);  
} else {  
    memcpy(pmemaddr, buf, cc);  
    pmem_msync(pmemaddr, cc);  
}
```

libpmemblk Examples

Source: <https://github.com/pmem/pmdk/blob/master/src/examples/libpmemblk>

```
/* store a block at index 5 */
strcpy(buf, "hello, world");
if (pmemblk_write(pbp, buf, 5) < 0) {
    perror("pmemblk_write");
    exit(1);
}

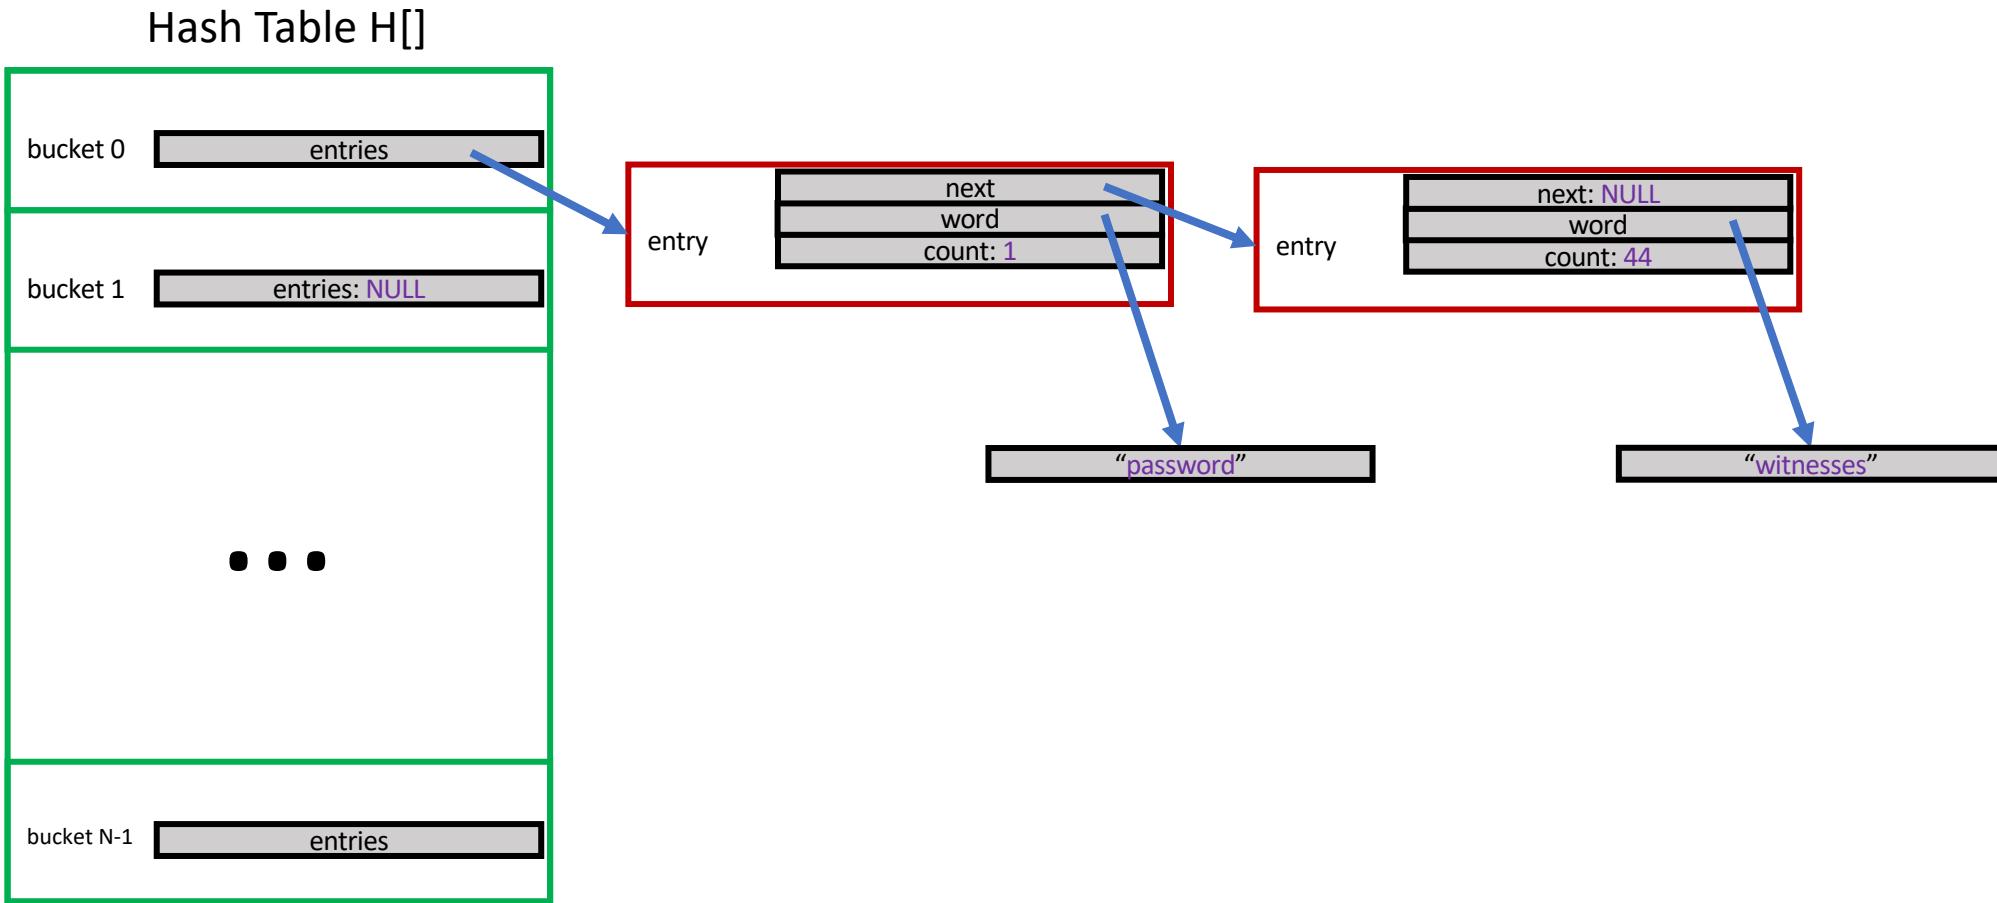
/* read the block at index 10 (reads as zeros initially) */
if (pmemblk_read(pbp, buf, 10) < 0) {
    perror("pmemblk_read");
    exit(1);
}

/* zero out the block at index 5 */
if (pmemblk_set_zero(pbp, 5) < 0) {
    perror("pmemblk_set_zero");
    exit(1);
}
```

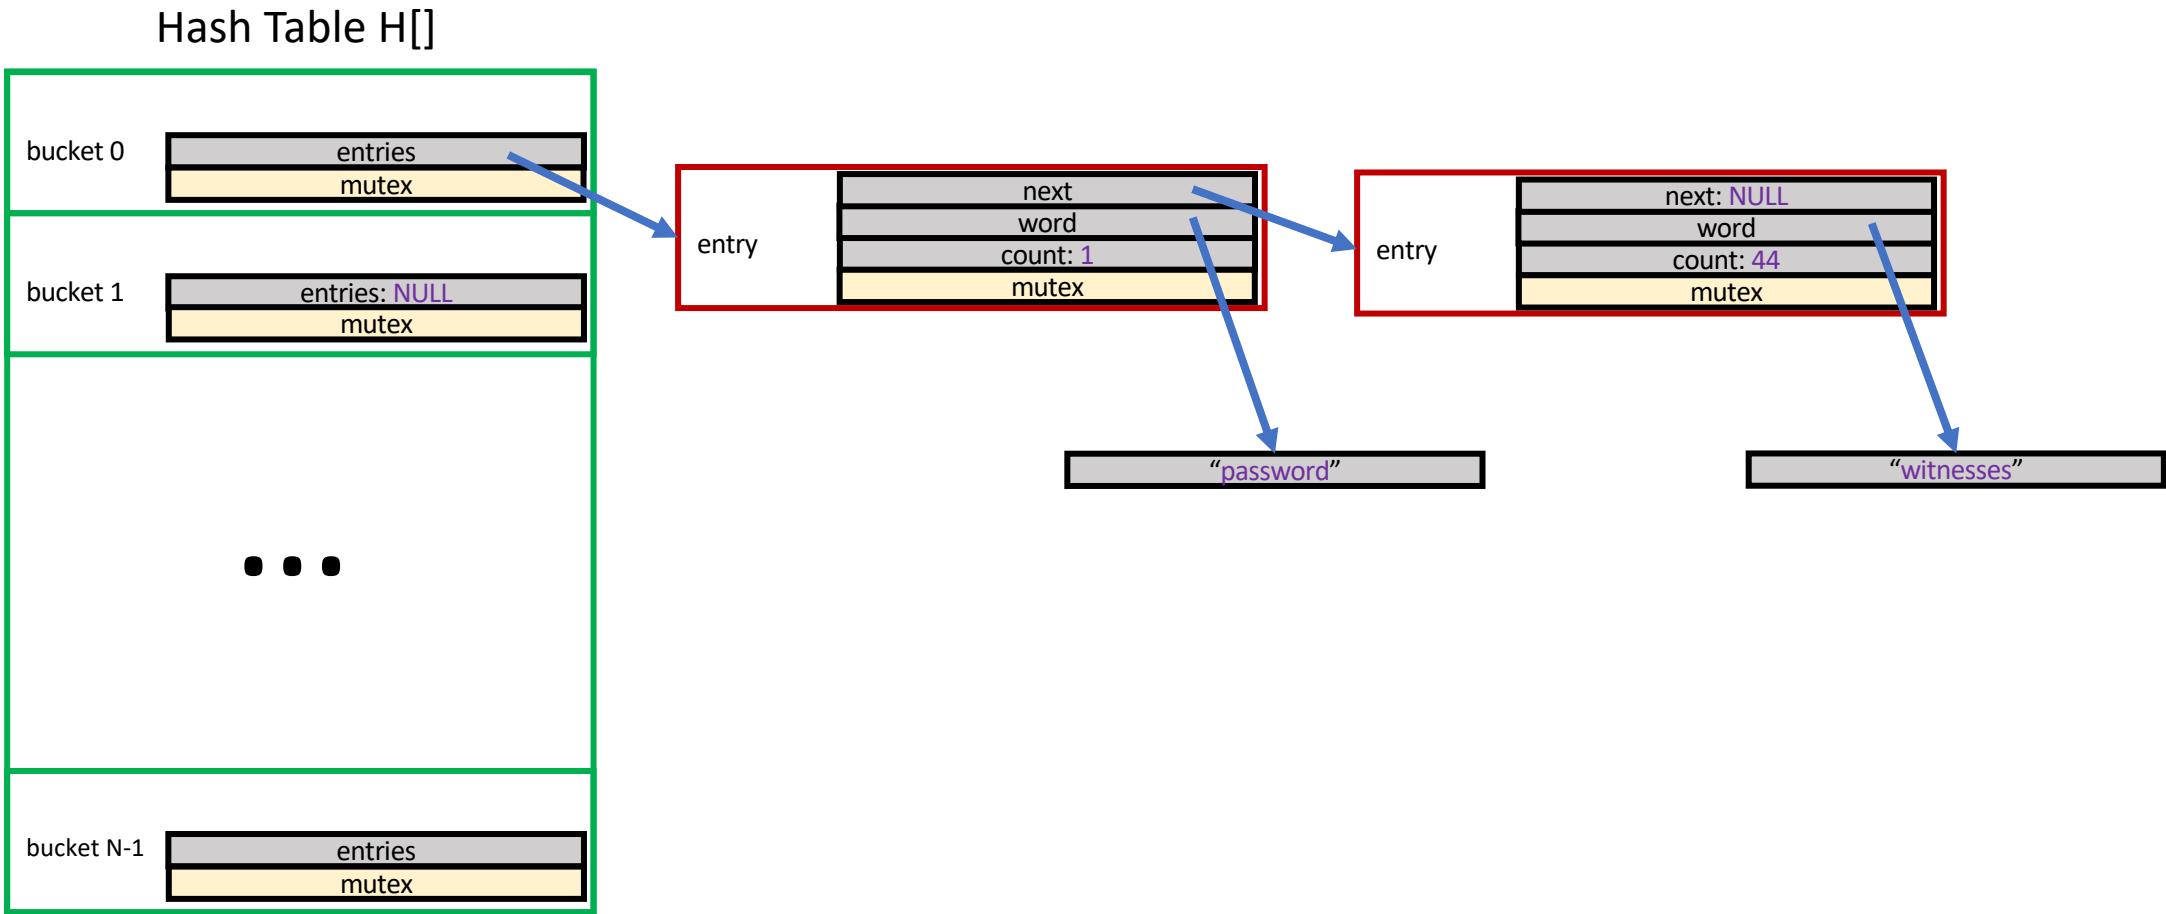
libpmemobj Examples

This is the most flexible of the PMDK libraries,
supporting general-purpose allocation & transactions

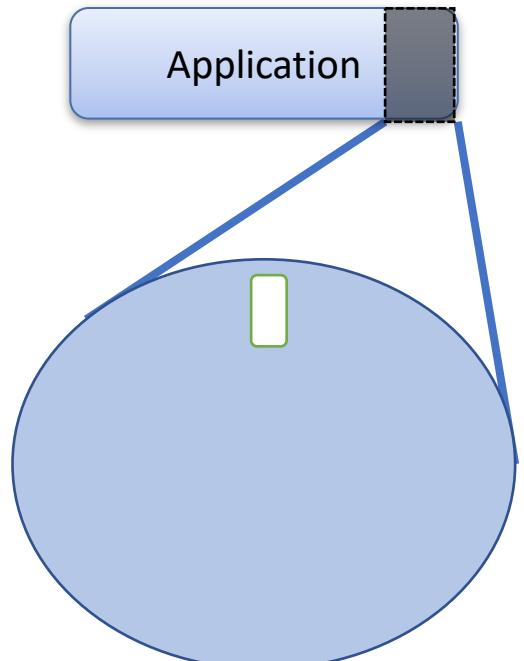
Simple C Example to Build On (no pmem yet)



Adding Multi-thread Support (no pmem yet)



The Root Object

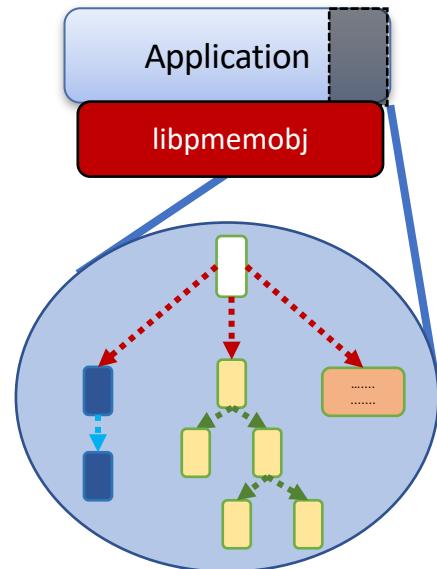


pmem pool “myfile”

root object:

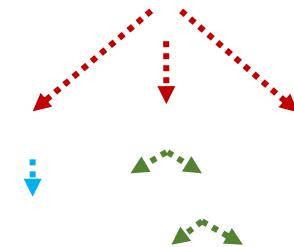
- assume it is always there
- created first time accessed
- initially zeroed

Using the Root Object

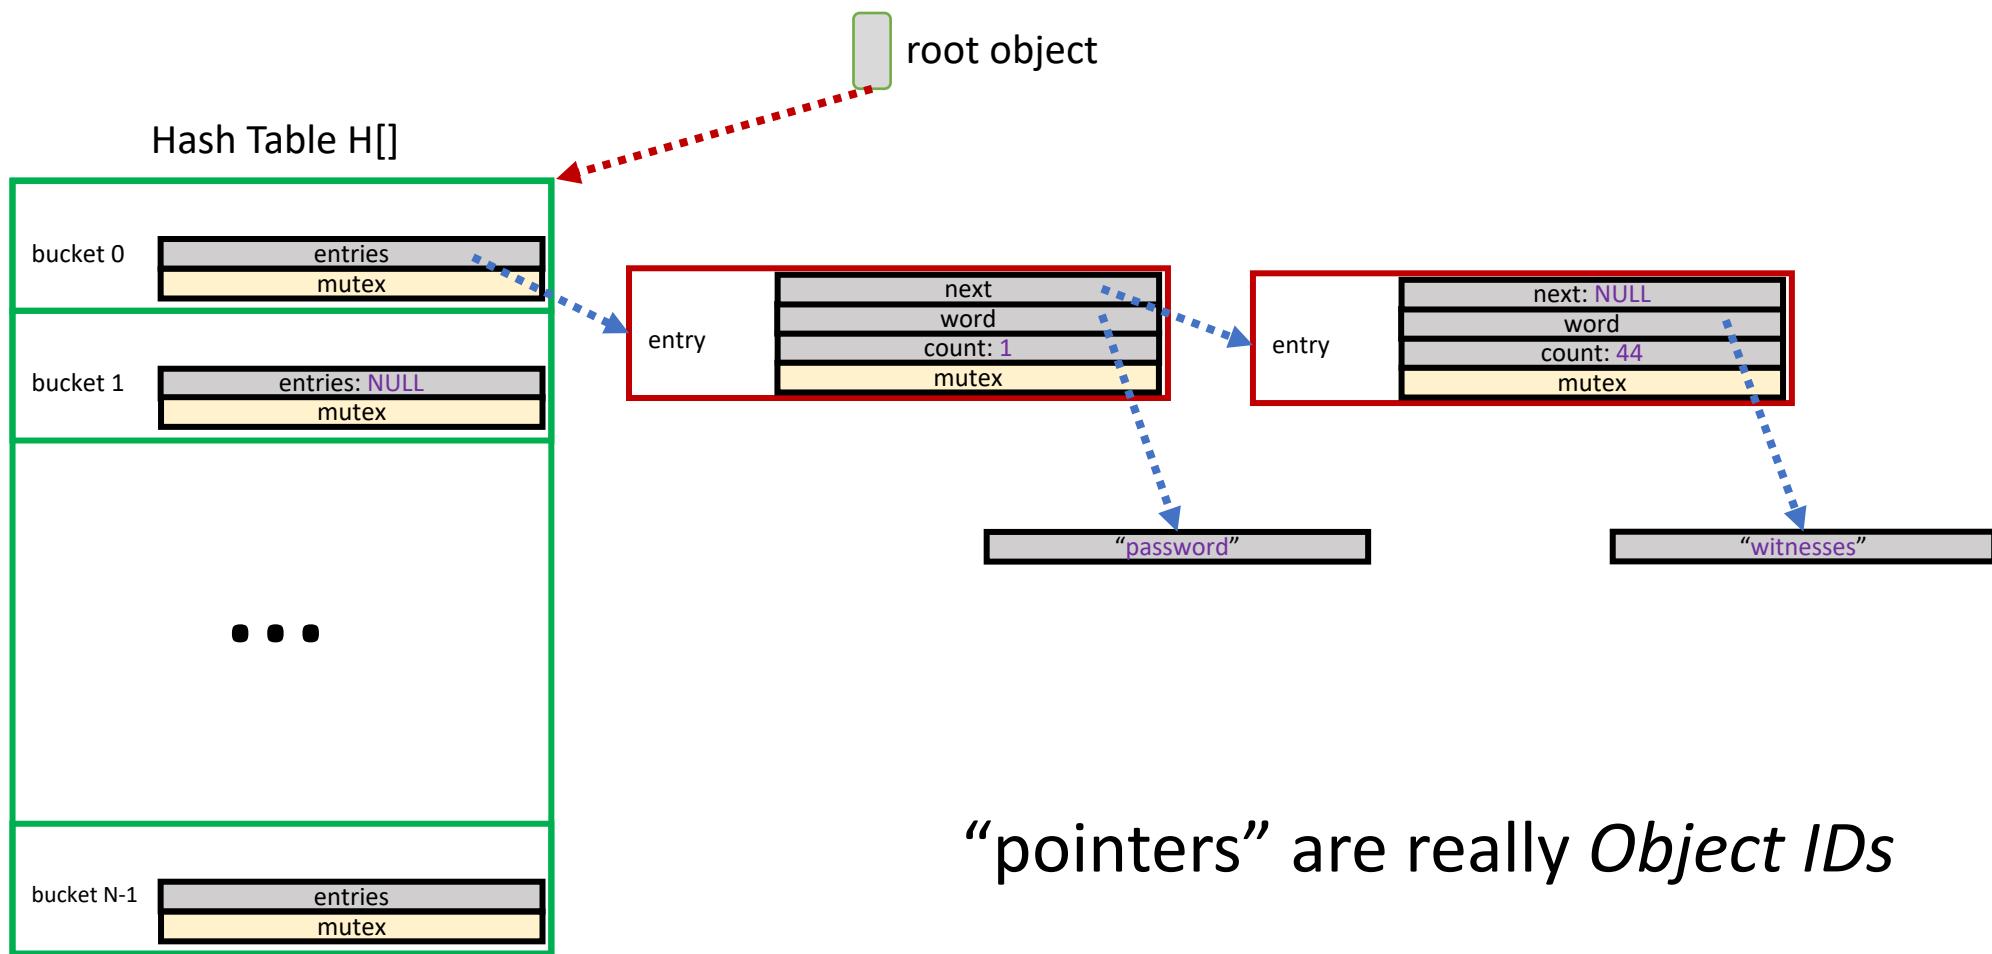


Link pmem data structures in pool off the root object to find them on each program run

“pointers” are really *Object IDs*



Moving the Example to pmem



C Programming with libpmemobj

Transaction Syntax

```
TX_BEGIN(Pop) {  
    /* the actual transaction code goes here... */  
} TX_ONCOMMIT {  
    /*  
     * optional - executed only if the above block  
     * successfully completes  
     */  
} TX_ONABORT {  
    /*  
     * optional - executed if starting the transaction fails  
     * or if transaction is aborted by an error or a call to  
     * pmemobj_tx_abort()  
     */  
} TX_FINALLY {  
    /*  
     * optional - if exists, it is executed after  
     * TX_ONCOMMIT or TX_ONABORT block  
     */  
} TX_END /* mandatory */
```

Properties of Transactions

Powerfail
Atomicity

Multi-Thread
Atomicity

```
TX_BEGIN_PARAM(Pop, TX_PARAM_MUTEX, &D_RW(ep)->mtx, TX_PARAM_NONE) {  
    TX_ADD(ep);  
    D_RW(ep)->count++;  
} TX_END
```

Caller must
instrument code
for undo logging

Persistent Memory Locks

- Want locks to live near the data they protect (i.e. inside structs)
- Does the state of locks get stored persistently?
 - Would have to flush to persistence when used
 - Would have to recover locked locks on start-up
 - Might be a different program accessing the file
 - Would run at pmem speeds
- PMEMmutex
 - Runs at DRAM speeds
 - Automatically initialized on pool open

C++ Programming with libpmemobj

C++ Queue Example: Declarations

```
/* entry in the queue */
struct pmem_entry {
    persistent_ptr<pmem_entry> next;
    p<uint64_t> value;
};
```

`persistent_ptr<T>`

Pointer is really a position-independent
Object ID in pmem.
Gets rid of need to use C macros like D_RW()

`p<T>`

Field is pmem-resident and needs to be
maintained persistently.
Gets rid of need to use C macros like TX_ADD()

C++ Queue Example: Transaction

```
void push(pool_base &pop, uint64_t value) {
    transaction::run(pop, [&] {
        auto n = make_persistent<pmem_entry>();
        n->value = value;
        n->next = nullptr;
        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Transactional
(including allocations & frees)

Links to More Information

More Developer Resources

- Find the PMDK (Persistent Memory Development Kit) at <http://pmem.io/pmdk/>
- Getting Started
 - Intel IDZ persistent memory- <https://software.intel.com/en-us/persistent-memory>
 - Entry into overall architecture - <http://pmem.io/2014/08/27/crawl-walk-run.html>
 - Emulate persistent memory - <http://pmem.io/2016/02/22/pm-emulation.html>
- Linux Resources
 - Linux Community Pmem Wiki - <https://nvdimm.wiki.kernel.org/>
 - Pmem enabling in SUSE Linux Enterprise 12 SP2 - <https://www.suse.com/communities/blog/nvdimm-enabling-suse-linux-enterprise-12-service-pack-2/>
- Windows Resources
 - Using Byte-Addressable Storage in Windows Server 2016 -<https://channel9.msdn.com/Events/Build/2016/P470>
 - Accelerating SQL Server 2016 using Pmem - <https://channel9.msdn.com/Shows/Data-Exposed/SQL-Server-2016-and-Windows-Server-2016-SCM--FAST>
- Other Resources
 - SNIA Persistent Memory Summit 2018 - <https://www.snia.org/pm-summit>
 - Intel manageability tools for Pmem - <https://01.org/ixpdimm-sw/>

Intel Developer Support & Tools

- **PMDK Tools**

- Valgrind plugin: pmemcheck
- Debug mode, tracing, pmembench, pmreorder

pmem.io

- **New features to support Intel® Optane™ DC persistent memory**

- Intel® VTune™ Amplifier – Performance Analysis
- Intel® Inspector – Persistence Inspector finds missing cache flushes & more
- Free downloads available

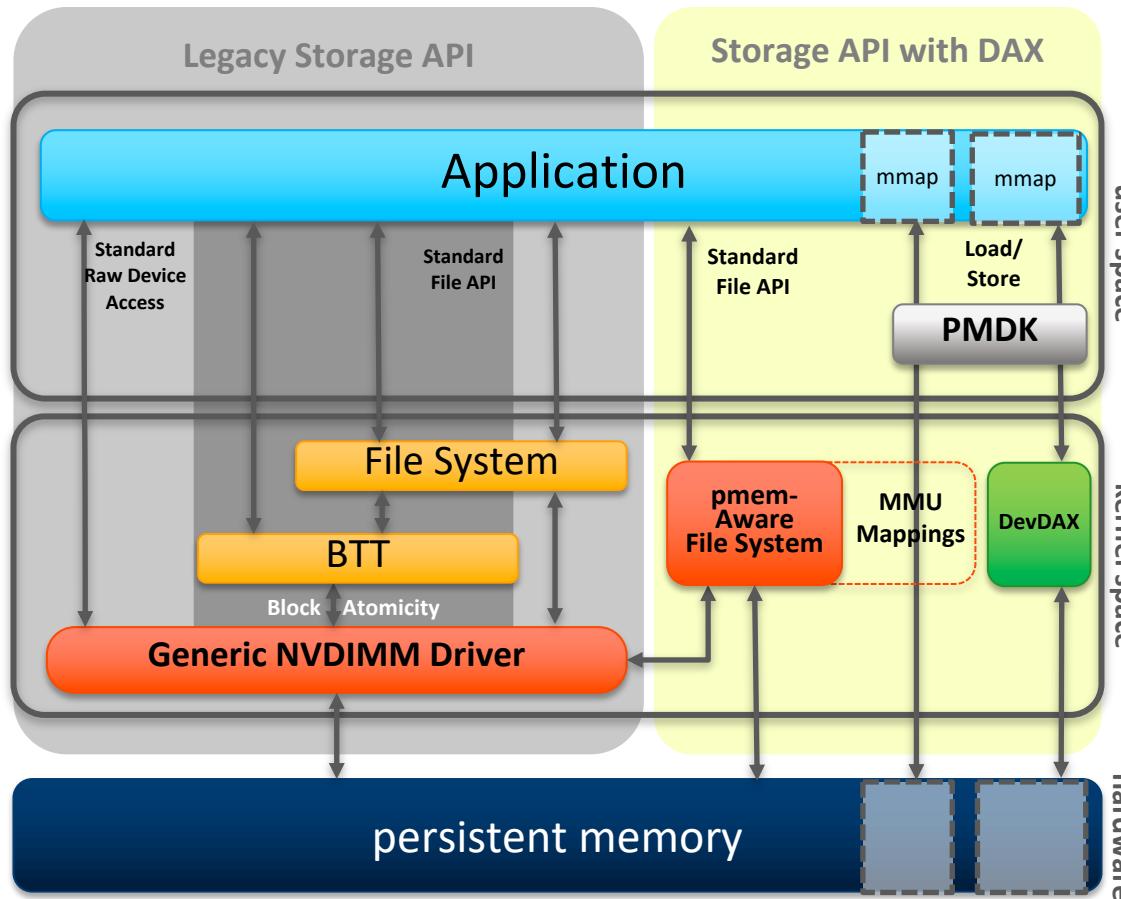
software.intel.com/pmem

Hack, hack, hack...

Possible ways to access persistent memory

- No Code Changes Required
- Operates in Blocks like SSD/HDD
 - Traditional read/write
 - Works with Existing File Systems
 - Atomicity at block level
 - Block size configurable
 - 4K, 512B*
- NVDIMM Driver required
 - Support starting Kernel 4.2
 - Configured as Boot Device
 - Higher Endurance than Enterprise SSDs
 - High Performance Block Storage
 - Low Latency, higher BW, High IOPs

*Requires Linux



- Code changes may be required*
- Bypasses file system page cache
- Requires DAX enabled file system
 - XFS, EXT4, NTFS
- No Kernel Code or interrupts
- No interrupts
- Fastest IO path possible

* Code changes required for load/store direct access if the application does not already support this.